


0



From Zero to Pipeline

How to build a go-to-market engine that turns positioning into pipeline — written for founders who can ship a product but have never had to sell one.

POSITIONING

ICP

FOUNDER-LED SALES

PRICING

90-DAY PLAN

FROM ZERO TO PIPELINE

The Modern GTM Playbook for B2B SaaS

How to build a go-to-market engine that turns positioning into pipeline.

Sunny Hooda

Founder, Great Hornbills · The Martech Guy

First edition · 2026

© 2026 Sunny Hooda. All rights reserved.

This book contains frameworks and benchmark data compiled from public research and industry reports. Every external statistic is attributed in the Sources appendix. Nothing here is investment, legal or financial advice. Your numbers will be your own.

Who this book is for

You built something. It works. People who try it usually like it. And almost nobody is buying it.

That gap is not a product problem and it is not a talent problem. It is a missing system. You have spent months building a machine that makes value, and roughly zero time building the machine that finds people who want it, convinces them, and gets money out of them. Those are two different machines. Nobody hands you the second one when you ship.

This book is the second machine.

You, specifically

I wrote this for a particular kind of founder, and it will feel uncannily well-aimed if that's you:

- **You're technical, or technical-adjacent.** You can build. You may have built the whole thing yourself, or with AI doing most of the typing, in weeks rather than years.
- **You're small.** One person. Three. Maybe ten. There is no marketing team. There is no sales team. There is you, and a to-do list that has "figure out marketing" on it in a slightly desperate handwriting.
- **You're pre-Series A.** Somewhere between your first paying customer and roughly \$1-2M ARR. Or at zero, and honest about it.
- **You sell to businesses.** Someone at a company has to say yes and expense it. That's B2B, and it behaves differently from consumer.
- **You have never done this before.** Not "I'm rusty." Never. You don't know what an ICP is, why anyone would care about positioning, or what people mean by "pipeline."

If you're at \$10M ARR with a VP of Sales, you already have an engine and you're tuning it. Read the benchmark appendix and skip the rest. If you're selling to consumers, most of the mechanics here don't transfer. If you're a marketer who already knows the frameworks, you'll find the 2026 data useful and the rest familiar.

What this book assumes

That you're smart, that your time is expensive, and that you don't want a lecture. Every chapter ends with something you actually do. There are fourteen worksheets. Fill them in and you will finish this book with a written go-to-market plan, not a feeling of vague inspiration.

It also assumes you'd rather be told the truth than flattered. Some of what's in here is going to be uncomfortable — that your product might not be as differentiated as you think, that your first thirty sales calls are going to be bad, that you personally have to do the selling for longer than you'd like. All of that is fixable. None of it is optional.

THE ONE SENTENCE VERSION

Pick one kind of customer, learn to say one true thing that makes them lean in, take that sentence to them through two or three deliberate channels, sell it yourself until it repeats, then and only then hire someone to do it at volume.

Everything else in these pages is detail on how to do that without wasting a year.

How to read it

Front to back, once, quickly — a few hours. Don't fill anything in on the first pass. Then go back to Part I and actually do the work, chapter by chapter. The order matters: positioning feeds the ICP, the ICP feeds the message, the message feeds every channel, and the channels feed the numbers. Skip to Chapter 9 looking for outbound tactics and you'll be sending well-written emails about a product nobody has been given a reason to care about.

Chapter 12 is a week-by-week plan that sequences the whole book into ninety days. If you only ever act on one page, make it that one.

What's in here

PART I · FOUNDATIONS — THE REASON NOBODY IS BUYING

- 01 The Build Trap
Why good products don't sell themselves, and what changed in B2B buying
 - 02 Positioning
The highest-leverage hour of your quarter, and the 10-step process to spend it well
 - 03 The ICP That Actually Filters
Five layers, a scoring rubric, and your first 100-account list
 - 04 Message
Turning a position into words that move money — in four lengths
-

PART II · PROOF — DOES ANYONE ACTUALLY WANT THIS

- 05 Product-Market Fit Is a Number
How to measure it, and the rule about not scaling before you have it
 - 06 Founder-Led Sales
Your first ten customers, the discovery call, the demo, the close
 - 07 Pricing and Packaging
Value metrics, the 2026 model mix, and how to price with zero data
-

PART III · THE ENGINE — MAKING IT REPEATABLE

- 08 Choosing Your Motion
Self-serve, sales-led, or hybrid — decided by ACV, not by taste
 - 09 The Five Doors
Where pipeline comes from, how to test channels, and how to pick two
 - 10 The Numbers
Your whole GTM model on one page, plus the 2026 benchmarks to judge it against
 - 11 AI as Leverage
What the data says AI actually improves in GTM — and where it quietly hurts
-

PART IV · SEQUENCE — WHAT TO DO ON MONDAY

- 12 The First 90 Days
A week-by-week operating plan that installs the whole engine
-

- 13** **Scaling Past Yourself**
When to hire, who first, and what breaks at every stage
-

APPENDICES

- A** **The Worksheets**
All fourteen, in order, ready to fill in
-
- B** **Benchmark Tables**
2026 conversion rates, cycles, costs, win rates, retention
-
- C** **Glossary**
Every GTM term in this book, defined for someone who has never sold anything
-
- D** **Sources & Further Reading**
Where the numbers came from
-

A NOTE ON THE NUMBERS IN THIS BOOK

Benchmarks are useful and dangerous. Useful because they tell you whether a 22% demo-to-close rate is good or terrible. Dangerous because they invite you to manage toward someone else's average instead of your own reality. Where I quote a figure, I name the source and the sample. Use them as a sanity check on your own data, never as a substitute for it.

PART ONE

Foundations

Almost every “we can’t get customers” problem is actually one of four problems wearing a disguise. You don’t know precisely what you are. You don’t know precisely who it’s for. You can’t say either of those things in a sentence. Or you’ve never asked the market whether it agrees.

This part fixes all four. It’s the least fun section of the book and by a wide margin the most valuable. Everything downstream — every email, every ad, every sales call, every landing page — either inherits the clarity you build here or inherits the confusion.

-
- | | |
|----|--|
| 01 | The Build Trap — what actually changed in B2B buying, and the six-part engine you’re missing |
|----|--|
-
- | | |
|----|---|
| 02 | Positioning — deciding what you’re the best at, for someone who cares |
|----|---|
-
- | | |
|----|--|
| 03 | The ICP That Actually Filters — five layers, a score, and a list of 100 real companies |
|----|--|
-
- | | |
|----|---|
| 04 | Message — the same truth at four lengths, from one line to one deck |
|----|---|

The Build Trap

There is a specific week that happens to almost every technical founder. The product works. You've fixed the embarrassing bugs. You put it live, told a few people, posted about it once. Twelve people signed up. Two of them logged in twice. Nobody paid. And you sit there thinking: I must need more features.

You almost certainly don't. What you need is a different kind of work than the kind you're good at, and the reason it feels so bad is that it doesn't produce anything you can look at when you're done.

Building has a wonderful property: effort converts to output reliably. You write code, the code runs, the thing exists. Go-to-market does not behave like this. You can do everything right for six weeks and see nothing. You can also do one thing right and have it pay you for two years. It's lumpy, delayed, and mostly invisible. Which is exactly why founders who are good at building avoid it, and then blame the product.

The uncomfortable arithmetic

Here's the part that most people building software in 2026 haven't sat with properly. The cost of building has collapsed. The cost of being noticed has not.

Five years ago, shipping a credible B2B SaaS product took a team, a year, and a seed round. Now a competent developer with AI assistance can put a working product in front of users in a few weeks. That is genuinely wonderful and it has a consequence nobody likes: your product being good is no longer rare. The scarce resource moved. It used to be engineering. It's now attention and trust.

Look at what happens to the companies that don't get this. Analysis of AI-layer startups puts failure rates for thin application-layer products at brutal levels — the large majority never generate meaningful revenue, and only a small single-digit percentage cross even \$10K MRR.¹ Not because the code was bad. Because a product that can be rebuilt in a weekend has to win somewhere other than the code, and most of them never tried.

Your product is a claim. Go-to-market is everything you do to make that claim credible to a stranger who has no reason to believe you.

What changed on the buyer's side

Meanwhile the person you're selling to got harder to reach, in four specific ways. These aren't vibes. They're measured.

1. It isn't one person anymore

Gartner's research puts the average B2B buying committee at around eleven people.² Eleven. Your champion — the person who loves your product — is one voice among ten others who have never seen a demo and have opinions about vendor risk, security review, and whether this overlaps with something they already pay for. More people in the room correlates with a large drop in the ability to reach a decision at all. Most deals don't die because you lost to a competitor. They die because a group of people failed to agree, and doing nothing was the safe outcome.

2. They barely talk to you

Gartner also finds buyers spend roughly 17% of their total purchase-evaluation time meeting with potential suppliers — and that sliver is split across every vendor they're considering.² If three vendors are in play, you get something like 5% of their attention. The other 95% is them reading, asking peers, watching videos, poking at your docs, and forming an opinion about you with no involvement from you whatsoever.

What you put in the world when you're not in the room is doing most of your selling.

3. Almost nobody is in the market right now

The 95:5 rule, from Professor John Dawes at the Ehrenberg-Bass Institute and published with the LinkedIn B2B Institute, says roughly 5% of your potential buyers are actively in a buying process at any given moment.³ The other 95% are not going to buy from you this quarter no matter how good your email is.

This single fact reorganises how you should think about marketing. If you only ever speak to people who are ready to buy, you're fishing in 5% of the pond and competing with everyone else fishing there. The companies that win the other 95% are the ones already sitting in the buyer's head when the need finally shows up — which takes time, and starts long before anyone raises a hand.

4. An AI is now the first thing they ask

This one is new and it's moving fast. Buyers increasingly start vendor research inside ChatGPT, Claude, Perplexity or Gemini rather than Google — Forrester's 2026 buying research puts generative AI at the top of buyer research interactions, ahead of both search and peer referral.⁴ When someone types "best tools for X for a company like mine," the three or four names that come back are the shortlist. Everything else vanishes before a human ever sees it.

And ranking on Google does not get you into that answer. One analysis of B2B SaaS buyer queries found that of pages ranking in the top three Google positions, only about 12% were actually cited by the major assistants.⁵ Different game, different rules. We'll come back to this properly in Chapter 9.

PUT TOGETHER

Eleven people have to agree. You get 17% of their evaluation time, split with competitors. Only 5% of them are even in the market this quarter. And their first research step happens inside a model you don't control. That is the environment. "Build a great product and tell a few people" is not a strategy that survives contact with it.

Distribution is the moat now

There's a version of this book that would tell you to go find a defensible technical moat. It's decent advice and mostly unactionable for a three-person team. The realistic version is this: for early-stage B2B software in 2026, your durable advantages are usually one of four things, and only one of them is code.

- **A workflow you become the system of record for.** Once you hold the data and the daily habit, ripping you out costs more than keeping you.
- **Proprietary data that compounds through use.** Every customer makes the product better for the next one in a way a competitor can't copy by writing the same features.
- **A vertical too narrow for a big player to justify.** Horizontal tools won't build for the specific ugly reality of a specific industry. You can.
- **Distribution.** An audience that trusts you, a channel you understand better than anyone else, a partner ecosystem, or a brand that comes to mind first.

Notice that three of the four are go-to-market decisions, not engineering ones. And distribution is the only one available to you starting this afternoon, with no funding and no data advantage. That's why it's the focus of this book.

The engine you're missing

When people say "go-to-market" they're usually gesturing vaguely at "sales and marketing." That's too fuzzy to act on. Here's the version I use, because it decomposes into six things you can actually build, in order, each one feeding the next.

FIGURE 1 — THE PIPELINE ENGINE

PART	THE QUESTION IT ANSWERS	WHAT BREAKS WITHOUT IT
1. Position	What are we, what do we replace, and what are we best at?	Every asset you write sounds like every competitor. Buyers can't tell you apart, so they pick on price or don't pick.
2. Audience	Who exactly, and how do I recognise them from the outside?	You market to "companies." Nothing lands. Your best channel looks like it failed because you pointed it at the wrong people.
3. Message	What do I say, in one line and in ten minutes?	Demos wander. Homepages list features. Emails get ignored. The value is real but invisible.
4. Offer	What are they buying, in what shape, at what price?	Deals stall at "send me pricing." You leave money on the table, or price yourself out of the segment that would have loved you.
5. Motion & channels	How do they find us, and how do they buy?	You try seven channels badly instead of two well. Cost per customer never comes down.
6. Model	How many of what, at what rate, to hit the number?	You can't tell whether things are working or you're just early. You optimise the wrong thing for two quarters.

Read that table again as a diagnostic. Whatever is going wrong in your business right now, it's in one of those six rows, and it's almost always the highest row that's broken. Fixing channel selection when your positioning is mush is repainting a car with no engine.

The four ways early founders lose

I want to name these plainly, because you'll recognise at least one and recognising it is most of the fix.

The feature homepage

Your site says what the product *does*. "AI-powered workflow automation for modern teams." Real-time sync. Enterprise-grade security. A grid of nine icons. A visitor reads it and learns nothing about whether it's for them or what changes in their week if they buy. This is the single most common failure and the cheapest to fix.

The everyone ICP

"Anyone with a support team could use this." Technically true. Commercially fatal. Broad targeting makes every message generic, every channel expensive, and every sales call a discovery mission from scratch. Narrowing is the fastest growth lever available to a small team and it feels, every single time, like giving up revenue.

The demo tour

Someone books a call. You are thrilled. You show them the product for 35 minutes, they say "this looks really interesting, let me discuss internally," and you log it as a good call. It was not a good call. Nobody named a decision owner, a success metric, a budget, or a next step. You got politeness and mistook it for progress.⁶

Channel roulette

Two weeks of cold email. Then some LinkedIn posts. Then a bit of Google Ads. Then a Product Hunt launch. Nothing gets enough time or volume to produce a signal, so everything "doesn't work," and you conclude the market is hard. Channels need a defined test with a defined budget and a defined kill criterion, which is Chapter 9.

THE REFRAME THAT MATTERS

Stop thinking of go-to-market as promotion — the thing you bolt on after the product is finished. Think of it as a second product. It has a spec, components, a build order, and a way of knowing whether it works. You already know how to build products. Build this one.

What "pipeline" actually means

Since it's on the cover, let's define it. **Pipeline** is the total value of the real opportunities currently in progress — people who have shown genuine intent, who you've qualified, and who have a plausible path to paying you. Not signups. Not newsletter subscribers. Not "people who liked my post." Opportunities, with a number attached and a date.

The reason pipeline is the right unit is that it's the earliest thing you can measure that reliably predicts revenue. Revenue tells you about a decision someone made weeks ago. Pipeline tells you about the decision they're making now. When you can see pipeline, you can forecast; when you can forecast, you can decide where to spend; when you can decide where to spend, you have a business rather than a hobby with invoices.

Getting from zero to a pipeline you can see and trust is the entire job of the next eleven chapters.

WORKSHEET 1

The honest diagnostic

Answer each out loud, in one sentence, without opening a browser tab. If you hesitate for more than five seconds, mark it as a no. Be brutal — the value of this page is entirely in the parts you get wrong.

- ☐ I can name the three things a prospect would realistically do instead of buying my product, including "nothing."
- ☐ I can name one capability I have that those alternatives cannot match, and I have proof, not an opinion.
- ☐ I can describe my best customer in a way that lets a stranger build a list of 100 more of them.
- ☐ I can say what my product does in one sentence with no words like "platform," "solution," "seamless," or "leverage."
- ☐ I know what problem my last five signups were actually trying to solve, in their words.
- ☐ I know why my last lost deal was lost — the real reason, not "budget."
- ☐ I know what my price is based on, other than a competitor's pricing page and a feeling.
- ☐ I know which single channel produced my last three customers.
- ☐ I know how many conversations it takes, on average, to produce one paying customer.
- ☐ I have written down what would make me stop doing my current main marketing activity.
- ☐ Someone other than me has repeated my value proposition back to me correctly.
- ☐ I could hand a new hire a document and have them run my sales process without me.

0–4 checked: You're pre-foundation. Do Part I in order, don't skip, don't buy any tools yet. **5–8:** You have instincts but nothing written down, which means nothing is repeatable and nothing can be delegated. Part I will take a week and will feel like it saved you six months. **9–12:** Your foundation is real. Go to Part III and fix the engine, and use Chapter 10 to find out which stage is actually leaking.

What to do before the next chapter

Two things, both small, both prerequisites for everything after.

- 1. Write down the last ten people or companies who engaged seriously with your product.** Paying, trialling, or just a real conversation. Name, company, size, what they were trying to solve, what happened. A spreadsheet is fine. This is your entire dataset and you're going to use it in Chapters 2, 3 and 5.
- 2. Book three conversations.** With customers if you have them, with people who churned if you don't, with people who said no if you don't have those either. Thirty minutes each. You're not selling. You're

going to ask what they were doing before you, and what they'd use instead if you disappeared.

Chapter 2 tells you exactly what to ask.

Nothing else. Don't redesign the site yet.

Positioning

Positioning is the single highest-leverage thing you will do all quarter, it takes about two days, and almost nobody does it deliberately. Instead they inherit it — from whatever they happened to build first, from the competitor they were reacting to, from the words that were in their head the day they wrote the homepage.

April Dunford, who wrote the standard work on this, defines positioning as **deliberately defining how you are the best at something that a defined market cares a lot about.**⁷ Read that slowly. Four constraints hiding in one sentence. *Deliberately* — a decision, not a default. *Best at something* — one thing, not a list. *Defined market* — someone specific. *Cares a lot* — enough to change what they're doing.

Miss any one of those and the whole thing collapses into “we’re a modern platform for growing teams,” which is the sound a company makes when it has decided nothing.

Why your instinct is wrong here

Founders position their product based on what they built. That’s the trap. You know every feature, you know why each one was hard, and you assume the value is self-evident from a list of them. But the buyer isn’t comparing your feature list to a blank page. They’re comparing it to *what they’re doing right now* — which is usually a spreadsheet, a person, a competitor, or nothing at all.

Positioning is not a description of your product. It’s a description of the gap between your product and the alternatives, from the point of view of someone who has to justify a purchase.

Nobody buys your product. They buy the difference between your product and what they’d do otherwise.

The five components

Everything in the process below is in service of filling in five boxes. Get these five right and the words write themselves.

FIGURE 2 — THE FIVE COMPONENTS OF POSITIONING

COMPONENT	WHAT IT MEANS
Competitive alternatives	What the customer would realistically do if you didn't exist. Almost always includes "nothing" and "a spreadsheet."
Unique attributes	The capabilities you have that the alternatives don't. Features, yes, but also data, integrations, model access, a team, a licence, a workflow.
Value	What those attributes let the customer <i>do</i> that matters. Not "faster reporting" — "the finance close happens on the 3rd instead of the 9th."
Who cares a lot	The characteristics of customers for whom that value is urgent rather than nice. This becomes your ICP in Chapter 3.
Market frame	The context you put yourself in so buyers instantly understand what you are and why your strengths matter.

The process, compressed for a small team

Dunford's full method is ten steps and assumes you have a marketing, sales and customer-success leader to convene.⁸ You don't. Here it is compressed into seven, in the order you'll actually run them. Give it two focused days. Do it with a whiteboard and your ten-customer spreadsheet from Chapter 1, not in your head.

Step 1 — Talk to the people who already love you

Before you decide anything, find out what's already true. You need three to eight conversations with people who use your product and would be annoyed if it disappeared. If you have no customers, talk to the people who came closest — the ones who trialled seriously, or the ones who said "not right now" rather than "no."

The mistake here is asking people to evaluate ideas. People are terrible at predicting their own behaviour and very good at being polite. Ask about the past instead. What they did, when, and why.

The positioning interview — 25 minutes, recorded, no pitching

Take me back to before you found us. What was going on that made you start looking?

What were you doing about it at the time? Walk me through it.

What else did you look at? Who else was on the list?

What made you pick us over those? Be blunt.

If we shut down tomorrow, what would you switch to? What would get worse?

How would you describe what we do to someone in your role at another company?

Who inside your company had to agree to this? What did they care about?

Answer 6 is the most valuable sentence in the call. Write it down verbatim. Answer 5 tells you your real competitive set. Answer 7 is your buying committee.

Two things you'll notice almost immediately. First, customers describe your product completely differently from how you do — usually more simply, usually in terms of an outcome you'd consider mundane. Second, the reason they bought is often not the feature you're proudest of. Both of those are gifts. Take the customer's words over yours every time.

Step 2 — List the true competitive alternatives

Not your competitors. The alternatives. There's a difference and it matters enormously.

Write down everything a prospect could realistically do instead of buying you. For most early B2B products the honest list looks like this:

- **Nothing.** They live with the problem. This is your biggest competitor and it wins most deals in the market. Never leave it off the list.
- **A spreadsheet, a Notion page, or a Slack channel.** Free, already paid for, understood by everyone, and surprisingly hard to beat.
- **A person.** An intern, an agency, an offshore contractor, an existing employee absorbing it into their week.
- **A general-purpose AI assistant.** New, real, and underrated as a competitor. If your product is a thin layer over what someone can prompt for directly, you will lose to this and you need to know it now.
- **The incumbent suite** they already pay for, which does 60% of what you do as a checkbox feature.
- **A direct competitor** — usually the smallest bucket, and the one founders overweight most.
- **Building it internally.** Especially if you sell to technical teams.

Then group them. You usually end up with two or three buckets that each require a different argument. The argument that beats “nothing” is about cost of inaction. The argument that beats a spreadsheet is about scale and error. The argument that beats a competitor is about a specific capability. If you try to write one message that beats all of them, you'll write something that beats none of them.

Step 3 — Isolate what you actually have that they don't

Now list every capability you have that the alternatives, as a group, do not. Be concrete and be honest — if a competitor has it too, it's table stakes, not a differentiator. Table stakes still need to exist; they just can't carry your positioning.

Look beyond features. Genuine differentiators for small companies often include:

- A data set nobody else has, or that gets better the more it's used
- An integration into the one system this segment actually lives in
- Deep knowledge of a narrow industry's ugly specifics — the regulation, the file format, the way the invoices actually come in
- Speed of implementation measured in hours instead of quarters
- A price point that makes you affordable to a segment the incumbents can't serve profitably
- You. The founder answering support within ten minutes is a real differentiator at this stage, and you should use it while you have it.

THE PROOF TEST

Every attribute you list needs evidence a sceptic would accept. “We’re easier to use” is a claim. “Median time from signup to first result is 11 minutes; the two products we replace take two weeks of onboarding” is proof. Third-party validation — a customer quote, a benchmark, a public number — is worth more than anything you say about yourself. If you can’t prove it, don’t build positioning on it.

Step 4 — Translate attributes into value

An attribute is what your product has. Value is what that lets a customer do. Run every attribute through the same question until you hit something a budget-holder would care about.

```
Attribute: "We auto-classify inbound support tickets."
  → so what?
Capability: "Tickets route to the right person without a human triaging."
  → so what?
Outcome:   "First-response time drops from 6 hours to 40 minutes."
  → so what?
Business:  "You stop losing renewals to slow support, and you don't
           hire the two extra tier-1 agents you'd budgeted for."
```

Stop when the next “so what” is obviously absurd. That last line is what goes on the homepage. The first line is what goes in the docs.

Then group the values into two or three **themes**. Most products have one dominant theme and one supporting one. If you have five, you haven’t finished thinking.

Step 5 — Decide who cares a lot

Here’s the step where founders flinch. Of the customers you could serve, which ones need this value *urgently*? Look at your ten-customer list and find the characteristics shared by the ones who bought fastest, paid most easily, complained least, and referred someone.

You’re looking for the thing that makes the pain acute rather than theoretical. Company size is rarely it on its own. Usually it’s a situation: a specific team structure, a recent event, a threshold they’ve just crossed, a regulation they’ve just come under.

A useful sharpening question: *who would be genuinely upset if we shut down?* Not disappointed. Upset. That’s your market.

Step 6 — Choose the market frame

This is the one that changes everything and gets the least attention. The market frame is the mental category you ask buyers to file you under. Change the frame and the exact same product goes from “expensive” to “cheap,” from “missing features” to “refreshingly focused.”

You have three options.

FRAME	WHAT IT MEANS	WHEN TO USE IT
-------	---------------	----------------

Head-to-head	You compete in an existing, well-understood category and claim to be better at it.	Only when you genuinely beat the leader on something buyers already rank highly. Rare for a small team. Expensive to defend.
Niche domination	Same category, but you're the obvious best choice for a specific slice of it.	Almost always the right answer pre-Series A. Cheap, credible, defensible, and it makes every channel work harder.
Category creation	You claim the existing categories don't describe you and name a new one.	Only if you have the budget and years to teach a market a new word. Seductive and usually a mistake at your stage.

Pick niche domination. I'll say it plainly: if you are under \$2M ARR with fewer than ten people, you should be the best-in-the-world at something narrow enough that a stranger can immediately tell whether they're in it or not. "Project management" is a category you'll lose. "Project management for structural engineering firms" is a category you can own, and owning it is what eventually lets you widen.

Step 7 — Layer a trend, carefully

Connecting yourself to a trend the market already believes in makes you easier to understand and easier to fund. It also makes you generic if the trend is the only thing you're saying.

Right now that trend is AI, and the phrase "AI-powered" has been so thoroughly used that it now carries close to zero information. It doesn't differentiate you and increasingly it invites the wrong question — *is this just a wrapper?* The application-layer failure data suggests buyers are right to ask.¹

The rule: a trend can be the adjective, never the noun. "AI" is not what you are. It's how you do the thing you are. If you strip the word AI out of your positioning and nothing is left, you don't have a position — you have a technology.

WORKSHEET 2

The positioning canvas

Fill this in after your interviews, not before. One page. If it takes more than one page you haven't decided yet.

1. COMPETITIVE ALTERNATIVES — WHAT THEY'D DO WITHOUT US

2. UNIQUE ATTRIBUTES — WHAT WE HAVE THAT THOSE DON'T (WITH PROOF)

3. VALUE THEMES — WHAT THOSE ATTRIBUTES LET CUSTOMERS DO

4. WHO CARES A LOT — THE CHARACTERISTICS THAT MAKE THIS URGENT

5. MARKET FRAME — THE CATEGORY WE WIN IN

6. THE STATEMENT

For _____ who _____, we are the _____ that _____. Unlike _____, we _____.

Three worked examples

These are illustrations, not case studies — invented products used to show the shape of the output. Notice how much more specific the “after” is, and how the market frame does most of the work.

Example A — a document tool

Before: “AI-powered document intelligence for modern teams.”

After: “Insurance brokers lose two days a week re-keying policy documents into their CRM. We read carrier PDFs — including the badly scanned ones — and write structured data straight into Applied Epic. Brokers using us cut renewal prep from two days to under an hour. Unlike general OCR tools, we already know what a schedule of insurance looks like.”

What changed: the frame moved from “document AI” (a crowded category where they’d be compared to enormous companies) to “renewal prep for insurance brokers” (a category with one obvious winner available).

Example B — a developer tool

Before: “The fastest way to monitor your infrastructure.”

After: “Teams under twenty engineers can’t justify a \$3k/month observability contract, so they run three free tools and find out about outages from customers. We’re one agent, one dashboard, flat \$99, live in ten minutes. Unlike the enterprise platforms, we don’t charge per host and we don’t need a solutions architect.”

What changed: the alternative was correctly identified as “a mess of free tools,” not the market leader. That reframes price as the feature.

Example C — a vertical workflow product

Before: “Workflow automation for healthcare.”

After: “Independent physiotherapy clinics lose 15–20% of revenue to insurance claim rejections they never chase. We check every claim against payer rules before it’s submitted and chase the rejections that get through. Clinics recover an average of £X per month. Unlike practice-management suites, we work alongside the system you already have — no migration.”

What changed: “no migration” is the whole deal. It removes the biggest reason for inaction, which was the real competitor all along.

How you'll know it's working

Positioning is not something you feel good about. It's something that produces observable changes within a few weeks:

- Prospects finish your sentences. They say “oh, so you're for X” and they're right.
- Discovery calls get shorter, because the unqualified people stop booking.
- Your win rate goes up even though your lead volume goes down.
- Somebody you've never met describes you correctly to somebody else.
- Writing content stops being agonising, because you know who you're talking to and what you're claiming.

And a warning: you will want to change it constantly. Don't. Positioning needs a couple of months of consistent exposure before the market can tell you anything. Change it when the evidence says it's wrong — when prospects consistently misunderstand you, when you keep winning deals for a reason you didn't plan — not when you're bored of it. You will be bored of it long before your market has heard it once.

DO THIS BEFORE CHAPTER 3

Book and run three positioning interviews using the script on page one of this chapter. Fill in Worksheet 2 with whatever you've got, even if it's rough. Chapter 3 turns box 4 — *who cares a lot* — into a list of a hundred real companies you can contact.

The ICP That Actually Filters

An ideal customer profile is not a description of who you'd like to sell to. It's a filter that lets you look at a list of ten thousand companies and confidently delete nine thousand nine hundred of them. If your ICP doesn't delete anyone, it isn't an ICP. It's a wish.

This is the chapter where you find out whether Chapter 2 worked. Positioning told you who cares a lot. Now you have to turn that into something operational — a set of observable characteristics that let you, or a tool, or a contractor in another timezone, build a list of real companies with real names.

Four words people use interchangeably and shouldn't

TERM	WHAT IT ACTUALLY IS	EXAMPLE
Market	Everyone who could conceivably buy. Useful for a pitch deck, useless for a Tuesday.	All companies with a support team
ICP	The <i>company</i> profile where you win most, fastest, and keep them longest. Firmographic and situational.	D2C e-commerce brands, 20–200 staff, on Shopify Plus, seasonal volume spikes, support team of 3–15
Persona	The <i>person</i> inside that company. Their role, their targets, their day.	Head of Customer Experience, measured on CSAT and first-response time, reports to COO
Buying committee	Everyone whose agreement is needed. Averages around eleven people in B2B. ²	Head of CX (champion), COO (economic buyer), IT (security review), Finance (procurement)

You need all four, but in that order, and the ICP is the one that changes what you do tomorrow morning.

The five layers

A good ICP has five layers. Most founders write the first one and stop, which is why their targeting doesn't work.

Layer 1 — Firmographic: what the company is

The structural facts. Industry, size (headcount and revenue), geography, business model, funding stage, growth rate. These are easy to find and easy to filter on, which is exactly why they're not enough on their own. Everyone can target "SaaS companies, 50–200 employees." That's not an advantage.

Layer 2 — Technographic: what they run

What's in their stack. This is more powerful than firmographics and most early founders ignore it. If your product plugs into Salesforce, then "uses Salesforce" is a harder filter than "is a mid-market company," and

it predicts both fit and speed of implementation. Companies with compatible stacks onboard faster and adopt more.

Sources for this are more accessible than people think: public job postings naming tools, integration marketplaces and partner directories, review-site profiles, engineering blog posts, and the technology detection built into most prospecting tools.

Layer 3 — Situational: what just happened to them

This is the layer that separates a list from a good list. A company that fits your firmographics perfectly but has no reason to change anything this quarter is a bad prospect. A slightly worse-fit company that just had the triggering event is a great one.

Triggers worth tracking, roughly in order of usefulness for early-stage B2B:

- **A relevant hire.** Someone was just given ownership of the problem you solve. New heads of department buy things in their first 120 days — it's how they establish that things are changing.
- **A hiring post** for the job your product makes unnecessary, or the job that would use your product.
- **Funding.** New money creates permission to spend and pressure to grow.
- **A threshold crossed.** Headcount, transaction volume, number of locations, entering a new market — whatever number makes your problem start hurting.
- **A public complaint.** Somebody in your target role posting about exactly the pain you solve. Rare, extremely high-converting.
- **A competitor relationship changing.** Price rises, an acquisition, a deprecated product, an outage that made the news.
- **A regulatory or compliance date** that applies to them and not to everyone.

The data supports leaning on this hard. Outbound that references a specific, verifiable trigger reports reply rates in a completely different band from generic cold outreach — industry benchmarks put cold reply rates around 3–4% while signal-led outbound reports 15–25%, with a corresponding 30–40% lower cost per qualified meeting.⁹ Treat those numbers as directional rather than gospel, but the direction is unambiguous and it matches what anyone who has done both will tell you.

Layer 4 — Behavioural: what they've done with you

Visited pricing three times. Opened a doc. Invited a colleague to a trial. Started and abandoned onboarding. Signed up with a work email from a target account. If you have any product at all, instrument this — it's the cheapest intent data you will ever own, and it's yours alone.

Layer 5 — Human: who signs and who suffers

For each account, three named roles:

- **The user** — who touches the product daily and whose life gets better.
- **The champion** — who will argue for you internally when you're not there. Often but not always the user. Look for someone whose personal performance metric moves because of you.
- **The economic buyer** — whose budget it comes out of and whose "no" ends it.

If those three are the same person, your sales cycle is short and your deal is small. If they're three different people in three different departments, your cycle is long and you need a champion who can sell

internally. Neither is better; they just require different books. Know which one you're in.

THE NEGATIVE ICP IS AS VALUABLE AS THE POSITIVE ONE

Write down who you will not sell to, and mean it. Companies below a size where the pain doesn't exist yet. Industries with procurement processes you can't survive. Anyone who needs a feature you have no intention of building. Anyone who has churned-customer characteristics.

Every founder can name a customer who cost them three months and left. Your negative ICP is the written version of that lesson, applied before it happens again.

Scoring: turning judgment into a number

Once you have the layers, weight them. This feels bureaucratic and takes twenty minutes, and its value is that it makes your targeting transferable — to a tool, a contractor, or your first hire.

FIGURE 3 — EXAMPLE ICP SCORING RUBRIC

CRITERION	WEIGHT	SCORING
Industry match	20	Exact target vertical = 20 · adjacent = 10 · other = 0
Size band (headcount)	15	In band = 15 · one band out = 7 · else 0
Runs required tech	20	Confirmed = 20 · likely = 10 · no = -10
Active trigger in last 90 days	25	Strong trigger = 25 · weak = 12 · none = 0
Named target persona exists	10	Found on LinkedIn = 10 · role absent = 0
Behavioural signal	10	Any product or site engagement = 10
Disqualifier present	-50	Negative-ICP characteristic = automatic exclusion

Then tier them. **Tier A (70+)** gets your personal attention — researched outreach, a custom angle, multiple touches across channels. **Tier B (40–69)** gets a lighter, templated sequence. **Tier C** goes in the newsletter and nowhere else. The point of tiering isn't tidiness; it's that your effort is finite and should be spent where it converts.

Building your first hundred

A hundred is the right number. Big enough to produce signal, small enough that you can research each one properly. Do not buy a list of ten thousand. You will send generic emails to all of them, burn your domain, and learn nothing.

Where the names come from

1. **Your own product data.** Everyone who signed up, trialled or enquired. The companies behind those email domains are your highest-intent source and they're free.
2. **Your customers' peer set.** For each existing customer, find ten companies that look like them. Same industry, same size, same tools. This is the single highest-yield list-building move at your stage.
3. **Job boards.** Search for the role that owns your problem, or the responsibilities in a JD that describe your product's job. Companies hiring for it have both the pain and a budget line.
4. **Review sites and marketplaces.** Who's writing reviews of the tool you replace or integrate with? Who's listed in the partner directory of a platform you plug into?
5. **Communities.** Slack groups, subreddits, Discords, LinkedIn groups where your persona complains in public. Read for six weeks before you post; you'll build a better ICP from that than from any database.
6. **Funding and news databases.** Filtered by your trigger, not by "recently funded" in general.
7. **Competitor footprints.** Public case studies, logo walls, event sponsor lists, conference attendee lists.

What to record for each

Keep it to one row and eight columns. Company, size, the evidence they fit, the trigger and its date, the named person, their role, the specific angle you'd open with, and the score. If you can't write a specific opening angle, the account isn't researched yet — and a hundred half-researched accounts are worth less than twenty properly researched ones.

WORKSHEET 3

ICP definition & account list

PART A — DEFINE

LAYER

OURS

Firmographic

Technographic

Triggers (top 3)

Behavioural signals

Champion role

Economic buyer role

Negative ICP — never sell
to

PART B — BUILD

- ☐ Listed every company that has ever signed up, trialled or enquired
- ☐ Found 10 look-alikes for each existing customer
- ☐ Ran two job-board searches for the owning role and saved the companies
- ☐ Identified 3 communities where the persona talks and joined all three
- ☐ Scored every account and tiered A / B / C
- ☐ Wrote a specific opening angle for all Tier A accounts
- ☐ Target: 100 scored accounts, of which at least 25 are Tier A

“But narrowing means less revenue”

It doesn't, and here's the arithmetic that shows why. It's the objection I get most, so it's worth doing properly.

Suppose you target broadly: 10,000 addressable companies, but your message fits none of them precisely. Say 1% respond, and of those you win 10%. That's 10 customers, and each one took a long, bespoke, exhausting sales cycle because you were learning their world from scratch every time.

Now narrow to 800 companies that share a situation, a stack and a trigger, and write a message about exactly that. Response goes to 6% because you're describing their actual Tuesday. Win rate goes to 25% because you have references who look exactly like them and you've heard every objection already. That's 12 customers — from a market 92% smaller.

10,000 accounts × 1% reply × 10% win = 10 customers	800 accounts × 6% reply × 25% win = 12 customers	1 of these two you can run as one person with no budget
---	--	--

Illustrative numbers, chosen to be conservative. The real effect is usually larger, because the narrow version also compounds: your case studies get more relevant, your content gets more findable, your referrals get denser, and your product roadmap stops being pulled in six directions by six unrelated customers.

Narrowing does not cap your ceiling. It sets your starting point. Every large software company you can name started absurdly narrow and widened once they'd won. You widen from strength, never from anxiety.

THE TEST FOR WHETHER YOUR ICP IS REAL

Hand your ICP definition to someone who doesn't work with you and ask them to build a list of twenty companies. If what comes back is roughly what you'd have built, your ICP is operational. If it's twenty random businesses, your ICP is still a sentence about how you feel.

Message

Positioning is a decision you make internally. Messaging is what comes out of your mouth and off your website. They're different jobs and founders collapse them, which is why so many companies have thought carefully about their strategy and still have a homepage that says "the all-in-one platform for modern teams."

The job of this chapter is narrow: take the six boxes from Worksheet 2 and turn them into a small set of sentences you'll reuse for the next year. Not copywriting flourish. Just the true thing, said in the clearest available words, at four different lengths.

The hierarchy

Everything you ever say about your company is one of five things. Write them once, in this order, because each level depends on the one above it.

FIGURE 4 — THE MESSAGE HIERARCHY

LEVEL	WHAT IT IS	WHERE IT SHOWS UP
Point of view	The thing you believe about this market that isn't universally agreed. Your argument, not your product.	Content, talks, the intro of a sales call
The one-liner	What you do, for whom, and what changes. Under 20 words.	Homepage headline, intros, ads, your own mouth
Value themes	The two or three outcomes you deliver, each with a plain-language explanation.	Homepage sections, deck, demo structure
Proof	Evidence a sceptic accepts for each theme.	Everywhere, immediately after every claim
Objections	The five reasons people don't buy, and your honest answer to each.	FAQ, sales calls, follow-up emails

Start with the point of view

Most B2B companies skip this and go straight to describing the product. It's a mistake, because a point of view is the only part of your messaging that people will voluntarily repeat, share, or argue with. Nobody forwards a feature list.

A point of view is a claim about how the world works that a reasonable person could disagree with. “Support teams should stop measuring ticket volume and start measuring how many customers had to contact them twice.” “Most sales problems are really onboarding problems.” “You don’t have a lead problem, you have a follow-up problem.”

The test: could a competent person in your industry read your point of view and say “no, I don’t think that’s right”? If not, it isn’t a point of view. It’s a platitude, and platitudes have no gravitational pull.

Your point of view should be the reason your product exists. If your product is the answer to a question nobody is asking out loud, your job is to make people ask it. That’s what content is for, which we’ll get to in Chapter 9.

The one-liner

You need one sentence that survives being said out loud at a party, typed into a LinkedIn bio, and read in 1.5 seconds at the top of a webpage. Three formulas that work, in rough order of how often they’re the right choice for a small company:

A. THE REPLACEMENT

“[Audience] use [old way] to [job]. We [new way], so [outcome].”

B. THE OUTCOME

“[Outcome that matters], for [specific audience].”

C. THE CONTRAST

“[Category] built for [specific situation], not [the default assumption].”

Worked through:

- A. “Insurance brokers re-key policy PDFs into their CRM by hand. We read the PDF and write the data, so renewal prep takes an hour instead of two days.”
- B. “Close the books on the 3rd. Built for finance teams at multi-entity startups.”
- C. “Observability for teams under twenty engineers — flat price, no solutions architect.”

Six rules for the sentence

1. **Use their words, not yours.** Go back to answer 6 in your positioning interviews — how customers describe you to a peer. Steal it. Their phrasing is better than yours because it was built for the ear of the person you’re trying to reach.
2. **Name the alternative.** “Instead of a spreadsheet” does more work than any adjective. It tells the reader which shelf you’re on.
3. **Be uncomfortably specific.** Numbers, timeframes, named systems, named roles. Specificity is the cheapest credibility available.
4. **No abstraction nouns.** Platform, solution, ecosystem, suite, framework, engine. Each one moves the reader further from a picture in their head.
5. **Say what changes, not what exists.** The reader doesn’t care that a dashboard exists. They care that they stop being surprised on the 28th.

6. **Cut every word that could apply to a competitor.** Read it back and delete any phrase that would still be true if a rival wrote it.

WORDS TO DELETE ON SIGHT

seamless · leverage · unlock · empower · revolutionary · game-changing · cutting-edge · next-generation · robust · scalable (unless you mean a specific number) · end-to-end · holistic · synergy · best-in-class · world-class · frictionless · supercharge · 10x · transform · streamline · all-in-one · modern teams · growing businesses · the future of anything

Every one of these is a word you reach for when you haven't decided what you mean. Delete it and you're forced to write the actual thing, which is always better.

The same truth at four lengths

You need your message in four sizes, and they must be the same message. The most common failure isn't bad copy — it's four different pitches in four different places, so the buyer who reads your ad, your homepage and your deck comes away with three impressions of three companies.

Five seconds — the headline

The one-liner, or half of it. Its only job is to make the right person keep reading and the wrong person leave. A headline that repels correctly is doing its job.

Thirty seconds — the paragraph

Problem, why the current way fails, what you do differently, one number as proof. Four sentences. This is your cold email body, your LinkedIn about section, your "so what do you do" answer.

The 30-second shape, filled in

"Brokers spend two days per renewal moving data out of carrier PDFs and into Epic. OCR tools don't help because they don't know what a schedule of insurance is, so someone still checks every field. We built the reader specifically for carrier documents — it maps straight into Epic and flags what it isn't sure about. The three brokerages using it now do renewal prep in under an hour."

Two minutes — the page

Homepage structure that works, in order: headline, sub-headline that names the audience, one proof element above the fold, the problem in the customer's words, your two or three value themes with a proof under each, how it works in three steps, objection handling, one call to action repeated. Nothing else. No feature grid of nine icons.

Twenty minutes — the deck and the demo

Same order as the page, expanded, with the demo slotted in as evidence rather than as the main event. Chapter 6 covers exactly how to run this. The key principle: the demo is a proof point in an argument, not a tour of a building.

Proof, ranked

Every claim needs evidence attached within a line or two of making it. Not all evidence is equal — here's the ranking, best to worst.

#	TYPE	WHY IT RANKS THERE
1	A named customer with a number	Specific, checkable, hard to fake. Worth ten of anything below.
2	A number from your own product data	"Median time to first result: 11 minutes." Yours, verifiable, and nobody else can copy it.
3	A live demo of the hard part	Showing beats telling. Show the thing that looks impossible.
4	Third-party research	Borrowed credibility. Good for framing the problem, weak for proving you solve it.
5	Logos with no context	Better than nothing, easy to dismiss, and dangerous if the logos are pilots.
6	Adjectives	"Powerful, intuitive, reliable." Zero evidential value. Everyone says them.

If you have no customers yet, you can still operate at level 2 and 3. Instrument your own product, measure something true, and publish it. "We ran this on 400 real documents from three brokerages and got 96.2% field accuracy; here are the 3.8% it got wrong and why" is stronger proof than most companies with revenue ever produce.

The five objections

Every B2B product faces roughly the same five objections, in the same order. Write your answer to each before you need it, because you will need all five within your first twenty sales calls and improvising them costs you deals.

OBJECTION	WHAT'S UNDERNEATH IT	WHAT ACTUALLY ANSWERS IT
"We already have something for this"	Change is expensive and the incumbent is good enough	Name the specific thing the incumbent doesn't do, and quantify what that costs them
"Now isn't the right time"	The pain isn't urgent, or the trigger hasn't happened	Cost of inaction with a number and a date. Or accept it and go dormant until the trigger fires
"It's too expensive"	Almost never about price. It's unproven value	Reframe against the cost of the alternative — the person, the hours, the failed renewals
"You're too small / too new"	Genuine risk aversion — what if you disappear	Reduce blast radius: short contract, easy export, a named human, a paid pilot with exit criteria
"I need to talk to the team"	You've been talking to a user, not a buyer	Give them the internal-sell kit and get on the next call. Never leave a champion to pitch alone

THE INTERNAL-SELL KIT

Your champion has to sell you to eleven people you'll never meet, in a meeting you're not invited to, using notes they wrote from memory. Make that easy. One page: the problem in their language, what you do, three bullet outcomes, what it costs, what happens in the first 30 days, and the two questions their CFO and their IT lead will ask with the answers already written. Send it unprompted. It is one of the highest-return documents you will ever write.

WORKSHEET 4

The messaging map

POINT OF VIEW — THE THING WE BELIEVE THAT OTHERS MIGHT NOT

ONE-LINER — UNDER 20 WORDS, NO BANNED WORDS

VALUE THEME 1 / PROOF

VALUE THEME 2 / PROOF

VALUE THEME 3 / PROOF (OPTIONAL — RESIST ADDING A FOURTH)

COST OF INACTION — WHAT IT COSTS THEM PER MONTH TO DO NOTHING

ANSWERS TO THE FIVE OBJECTIONS

Three tests before you ship it

1. **The five-second test.** Show your homepage to someone in your ICP for five seconds, take it away, and ask what the company does and who it's for. If they can't answer both, the headline has failed. Do this with five people; it costs nothing and it's the most reliable feedback loop in marketing.
2. **The stranger test.** Say your one-liner to someone who doesn't know your business. If their next sentence is a question about what you mean, rewrite. If it's "oh, so you do X for Y" — and X and Y are right — you're done.
3. **The call test.** On your next five sales calls, count how many times you have to improvise an explanation because your written message didn't cover it. Those gaps are your messaging backlog. The best copy you'll ever write is a transcript of you explaining your product well on a call at 4pm on a Thursday. Record your calls and mine them.

THE THING TO REMEMBER

Your message will feel repetitive to you roughly a thousand times before it registers with your market even once. Founders rewrite their positioning every six weeks out of boredom and then wonder why nobody knows what they do. Write it well, then say the same thing for a year.

PART TWO

Proof

You now have a position, an audience and a message. All three are hypotheses. This part is where you take them outside and find out whether the market agrees — by measuring whether anyone would miss you, by selling to twenty people yourself, and by asking them for money.

Nothing in Part III works until Part II is done. Scaling a motion you haven't personally proven is the most expensive mistake available to an early company, and it's the default behaviour of founders who find selling uncomfortable.

-
- 05 Product-Market Fit Is a Number — how to measure it, and what to do at each score
 - 06 Founder-Led Sales — the discovery call, the demo, the close, and the pipeline you keep
 - 07 Pricing and Packaging — value metrics, model choice, and pricing with no data
-

Product-Market Fit Is a Number

Product-market fit gets talked about like weather — something that arrives, that you sense, that people describe in retrospect with phrases like “you’ll just know.” That’s unhelpful when you’re two months in at 3am trying to decide whether to keep going or change something.

You can measure it. Not perfectly, but well enough to make decisions, and measuring it is the difference between iterating with intent and iterating with vibes.

The 40% test

Sean Ellis — the first marketer at Dropbox — benchmarked more than a hundred startups by asking their users a single question:

“How would you feel if you could no longer use this product?”

Very disappointed · Somewhat disappointed · Not disappointed

Companies where more than 40% answered “very disappointed” generally went on to grow. Companies below that threshold generally struggled, regardless of how much they spent on marketing.¹⁰

The question is well-designed and the design is the point. It doesn’t ask if people like your product — people are polite. It asks about loss, and people are honest about loss. It measures necessity rather than approval.

How to actually run it

- **Ask the right people.** Users who have experienced the core of your product at least twice in the last two weeks. Not everyone who ever signed up — including people who never got started measures your onboarding, not your fit.
- **Get enough responses.** Forty is workable, a hundred is comfortable. Below thirty, treat the number as a conversation-starter rather than a measurement.
- **Ask three follow-ups.** What type of person do you think would benefit most from this? What’s the main benefit you get? How could we improve it for you? These three open questions carry more information than the score itself.
- **Keep it repeatable.** Run it quarterly, same wording, same audience definition, so the trend means something.

The move that actually improves the score

Here's the part most people miss, and it's the whole reason to run the survey rather than just guess.

Rahul Vohra ran this at Superhuman and started at 22% — well below the bar.¹¹ Instead of treating that as a verdict on the product, he segmented. He looked only at the people who said “very disappointed,” worked out what they had in common, and re-scored within that segment. It came out at 33%. Then he built almost exclusively for that segment: doubled down on what they loved, and only fixed the blockers that the “somewhat disappointed” group named — the ones who were nearly converted. Over roughly three quarters, the score reached 58%.

The lesson generalises well beyond surveys. **An average PMF score across a mixed audience is meaningless.** You almost never have “40% fit.” You have 70% fit with one kind of customer and 10% with everyone else, averaging into a number that tells you to give up. Segment first. Then decide.

NOTICE WHAT THIS IS

It's Chapter 2 and 3 arriving through a different door. “Who are the people who'd be very disappointed, and what do they have in common?” is exactly the positioning question — *who cares a lot* — answered with data instead of interviews. If your PMF survey and your ICP disagree, believe the survey.

The four signals that matter more than the survey

The 40% test is a good instrument and a lagging one. These four are earlier and harder to fool.

1. Does retention flatten?

Plot the percentage of a cohort still active at week 1, 2, 4, 8, 12. Every product declines. The question is whether the curve *flattens* at some non-trivial level or keeps sliding toward zero. A curve that flattens at 35% means you have a real product for a real subset. A curve that reaches zero means you have a demo that people tried.

Flattening beats every other early signal, including revenue. You can buy signups. You cannot buy a flat retention curve.

2. Are people coming without being pushed?

What proportion of new signups arrived because someone told them about you, versus because you chased them? Word of mouth is the most honest metric in early-stage software because it costs the recommender social capital. Ask every new user how they heard about you, in a free-text box, and read the answers weekly.

3. Does usage frequency match the job's frequency?

If your product solves a daily problem, do people use it daily? Weekly problem, weekly? Mismatch here usually means they're using you for something smaller than you designed for — which is useful information and often a better business than the one you intended.

4. Do people pay without a fight?

Not “did they buy after a nine-week negotiation and a 40% discount.” Did they say a version of “yes, how do I pay?” Friction at the payment step is often the market telling you the value is real but smaller than your price, or real for the user but not yet real for the budget-holder.

FIGURE 5 — READING THE SIGNALS TOGETHER

STATE	WHAT YOU'D SEE	WHAT TO DO NEXT
No fit yet	Score under 25% in every segment. Retention slides to zero. No unprompted referrals.	Do not spend on distribution. Go back to Chapter 2, talk to twenty more people, be willing to change the product or the segment.
Fit with a segment	Under 40% overall, but 40%+ inside one identifiable group. Retention flattens for that group only.	Re-cut your ICP to that group and build for them exclusively for two quarters. This is the most common real state and it's a good one.
Fit	40%+ in your target segment, retention flattens, referrals appear, deals close without heroics.	Now scale. Part III. Add channels, add people, add spend — in that order.
False positive	Great early usage, sharp drop after week two. High trial numbers, low conversion.	You have a novelty product or an onboarding problem. Look at what activated users did differently in their first session.

A warning specific to AI products

If your product is AI-native, the standard signals lie to you more than usual, in a specific way: novelty produces a usage spike that looks exactly like enthusiasm and decays like a fashion. People try the impressive thing, tell a friend, and never return — not because it's bad, but because it wasn't attached to anything they have to do every week.

Three adjustments:

- **Measure week 4, not week 1.** Your week-1 numbers are measuring curiosity. Week 4 is measuring habit.
- **Separate “wow” from “work.”** Track the specific action that means someone integrated you into a real workflow — connected a data source, invited a colleague, ran it on production data. That number, not signups, is your real top of funnel.
- **Watch the buyer's clock, not yours.** Buyers are behaving as if today's best tool may be obsolete in six to twelve months, and they're contracting accordingly — contracts under one year jumped from 4% of new-logo deals in 2025 to 13% in 2026, while three-year deals fell.¹² Retention has to be earned continuously now. Fit that decays is not fit.

The rule about scaling

This is the most expensive lesson in early-stage go-to-market, so it gets its own box.

DON'T SCALE WHAT DOESN'T REPEAT

Before you hire a salesperson, buy ads, or run a hundred-account outbound campaign, you need a motion that has produced a customer more than once, in a way you can describe in writing. Not “we got some customers.” A repeatable path: this type of company, found this way, with this message, converts at roughly this rate, in roughly this many weeks.

Spending money to accelerate an unproven motion doesn't test it faster. It just makes the failure larger and later.

What you can legitimately do before you have fit

“Don't scale” is not “don't do anything.” Pre-fit, the correct activities are the ones that generate learning per unit of effort:

- Founder-led outbound to a narrow list — 30 highly researched accounts, not 3,000
- Ten customer conversations a month, forever, regardless of what else is happening
- Publishing your point of view, because it compounds and costs only time
- Manual, unscalable onboarding — doing it by hand teaches you what to automate
- Charging money from the first customer, because free users don't tell you the truth

What you should not do: hire an SDR, sign an annual tool contract, build a “marketing site,” run paid acquisition, or write a content calendar for six months out.

WORKSHEET 5

PMF scorecard

MEASURE	YOURS	BAR
“Very disappointed” %, all users		40%
“Very disappointed” %, best segment		40%+
What that segment has in common		
Week-4 retention (cohort)		Curve flattens
% of signups from word of mouth		Rising
Usage frequency vs job frequency		Matched
Deals closed without discounting		Most
Logo churn, last 90 days		Falling

Decision rule. Two or more rows below the bar: stay in Part II, keep selling by hand, keep changing the product or the segment. All rows at or above: you have permission to spend money.

Founder-Led Sales

You have to do this yourself. Not because it's noble, and not because you'll be good at it — you won't be, for a while. Because the first twenty sales are not a revenue activity. They're a research activity that happens to produce revenue, and you cannot outsource research about your own business.

Every experienced operator says the same thing about the handoff: carry it yourself until roughly \$500K–\$1.5M ARR, and the real trigger isn't a revenue number, it's repeatability.¹³ A founder at \$2M with chaotic conversion rates can't hire successfully. A founder at \$500K with clean, understood rates can. We'll come back to that in Chapter 13.

The reframe that makes this bearable

Most technical founders hate sales because they think it's persuasion — talking someone into something using confidence they don't feel. That version of sales is both unpleasant and ineffective in B2B, where the buyer has eleven colleagues and a security review and no interest in your charisma.

The version that works is diagnosis. You are trying to find out, as fast as possible, whether this person has a problem you can solve, whether it hurts enough to act on, and whether they can actually buy. Most of the time the answer is no, and finding that out in twelve minutes instead of six weeks is a win, not a failure.

A good sales call is one where you learn the truth quickly.
Sometimes the truth is a deal.

This reframe does something practical: it makes disqualifying feel good. Founders lose entire quarters to prospects who were never going to buy but were too polite to say so. If your job is diagnosis, “this isn't for you” is a successful outcome, and you'll say it much earlier.

The shape of the thing

Here's the whole process. Six stages, and you should be able to say which stage every live deal is in without thinking.

FIGURE 6 — THE FOUNDER SALES PROCESS

#	STAGE	EXIT CRITERIA — NOT VIBES	YOUR JOB
1	Target	Scored Tier A, named person, specific angle written	Research, 10 minutes max per account
2	Contact	They replied, or booked	Reach out in a way that references their situation
3	Discovery	Problem confirmed, cost quantified, next step booked	Ask, listen, write down their exact words
4	Prove	They've seen it solve <i>their</i> problem, not a generic one	Tailored demo, or a scoped pilot
5	Decide	Economic buyer engaged, paper process known, date agreed	Remove risk, arm the champion, hold the plan
6	Onboard	They hit first value inside the agreed window	Do it manually, personally, obsessively

Two rules about stages. First, a deal only advances when the buyer does something, never when you do something — sending a proposal is not progress, receiving feedback on it is. Second, if a deal has no scheduled next step with a date, it isn't in your pipeline. Move it out. Half of what founders call pipeline is hope with a company name attached.

Time: the unglamorous part

Two to three hours a day, every working day, ring-fenced. Outreach and research in one block, calls in another. If you're pre-fit and spending under a third of your time on customer-facing work, you're building too much.

Protect it the way you'd protect deploy windows. The reason founder-led sales fails is almost never a lack of skill. It's that a bug came in, and then a hiring call, and then it was Thursday.

The discovery call

This is the highest-value thirty minutes in your company. Get it right and everything downstream is easier — the demo writes itself, the objections are pre-empted, the proposal is a formality. Get it wrong and you will do a 35-minute product tour to someone who nods, says "this looks really interesting," and evaporates.

The one rule

Ask about the past, not the future. People are unreliable narrators of what they will do and reliable narrators of what they have done. "Would you use something like this?" gets you a yes that means nothing. "What did you do about it last month?" gets you the truth.

The structure

Discovery call — 30 minutes. You should talk for less than 30% of it.

Frame it. “I’ve got about thirty minutes. I’d like to spend most of it understanding how you handle X today, and if it makes sense I’ll show you what we do — but I’d rather not demo something that isn’t relevant. Fair?”

Context. “Tell me about your team and how X works today. Who touches it?”

The problem. “Where does that break down? — When did that last happen? — What did you do? — How often does that happen? — Who else does it affect?”

Quantify. “Roughly how many hours a week does that take? — What does it cost when it goes wrong? — Has anyone put a number on it internally?”

Alternatives and history. “What have you tried? — Why didn’t that work? — Is anyone else looking at this? — What happens if you do nothing for another six months?”

Process. “If you decided this was worth solving, what happens next inside your company? — Who else would need to be involved? — Is there budget for this, or would it need to be found?”

Show one thing. Not a tour. The single screen that addresses the thing they just described. “You mentioned the re-keying — here’s what that looks like with us.”

Next step, with a date. “Based on what you’ve said, I think the useful next step is X. Does Tuesday the 14th at 10 work, and can we get [name] on it?”

The four questions that do most of the work

1. **“Walk me through the last time that happened.”** Converts abstraction into a story with details you can sell against.
2. **“What did you do about it?”** Reveals the real competitive alternative, which is usually “nothing” or “Sarah stayed late.”
3. **“What does that cost you?”** If they can’t answer, help them build the number on the call. A problem without a number attached will never beat a budget process.
4. **“What happens if you don’t fix this?”** The single best predictor of whether a deal closes. If the honest answer is “nothing much,” you’re talking to someone who will enjoy the demo and never buy.

THINGS TO STOP DOING ON CALLS

Answering questions you weren’t asked. Founders fill silence with features. Let it be quiet.

Pitching before you’ve diagnosed. If you’re talking in the first ten minutes, you’re doing it wrong.

Accepting “this looks great” as progress. Follow it with: “what would need to be true for you to use this?”

Skipping the money question because it’s awkward. It gets more awkward later, not less.

Ending without a dated next step. “I’ll follow up next week” is where deals go to die.

Qualification without the enterprise theatre

The standard framework in B2B software is MEDDPICC — eight things you need to know before you can forecast a deal honestly.¹⁴ It’s built for six-figure enterprise deals and it’s overkill for a \$500/month

product. But the underlying questions scale down, and using a reduced version keeps you honest about which deals are real.

FULL MEDDPICC		WHAT YOU NEED TO KNOW AT YOUR STAGE
M	Metrics	The number that improves, and what it's worth. Always required.
E	Economic buyer	Who signs. Always required. If you haven't met them by stage 5, you don't have a deal.
D	Decision criteria	What they'll compare you on. Always required.
D	Decision process	Steps and dates. Required above ~\$10K ACV.
P	Paper process	Legal, security, procurement. Required above ~\$25K, or in regulated industries at any size.
I	Identified pain	Quantified, in their words. Always required.
C	Champion	Someone whose own metric improves. Always required.
C	Competition	Including "do nothing," which is usually who you're actually against.

For a small-ACV product, four are non-negotiable: **pain, metric, economic buyer, champion**. If you can't fill those four in on a deal, it isn't qualified, no matter how enthusiastic the call was.

The demo

The demo is not a product tour. A tour shows everything, which forces the buyer to work out which parts matter to them — work they will not do. A demo is an argument in three acts, and every screen you show is a piece of evidence for a claim you already agreed on in discovery.

FIGURE 7 — DEMO STRUCTURE

ACT	TIME	CONTENT
Recap	3 min	"Last time you told me X, Y and Z, and that Z costs you about 12 hours a week. Is that still right? Anything changed?" Get a yes. You are now demoing against an agreed problem.
The moment	5 min	Go straight to the single most valuable thing, with their data or a close facsimile of it. Not the login screen. Not the settings. The moment where the painful thing becomes easy.
The path	10 min	Now show how they get from today to that moment. Setup, integration, who does what in week one. This is where you kill "how long will this take."
The objections	7 min	Raise their likely objections yourself before they do — security, migration, the thing you don't do. Founders who name their weaknesses close more, because credibility transfers.
The step	5 min	Agree what happens next, who's on it, and when. In writing, same day.

Use their data if you can possibly get it. A demo on a sample dataset is a story. A demo on their actual documents, their actual accounts, their actual mess, is proof. Ask for it in the discovery call: "if you can

send me two of the ugliest ones, I'll run them before we talk again.”

Pilots and proofs of concept

For anything above a few thousand dollars a year, buyers increasingly want to try before they commit — and the data says this is now the strongest conversion path in B2B software. In 2026, free trial and proof-of-concept motions converted to paid at around 50% on average, against 28% for SQL-to-closed-won and 38% for demo-to-closed-won.¹⁵

50% Free trial / POC to paid	38% Demo to closed-won	28% SQL to closed-won	+14pt POC conversion, year on year
--	----------------------------------	---------------------------------	--

Source: ICONIQ survey of GTM executives at 150+ B2B software companies, January 2026.¹⁵

But that 50% is not what happens if you hand someone a login and hope. The operators in that data are explicit that the conversion comes from treating a pilot as a disciplined operating motion — heavy preparation before it starts, and active work alongside the customer during it.¹⁵ An unmanaged pilot is a way of losing slowly while feeling productive.

Rules for a pilot that converts

- **Written success criteria, agreed before it starts.** “By 14 March, we will have processed 200 of your documents at 95%+ accuracy and your team will have spent under two hours on setup.” Both sides sign off.
- **A fixed end date.** Open-ended pilots become free products. Two to four weeks for small deals, six at the outside.
- **Named people on both sides** and a scheduled mid-point check-in. Not “let me know how it goes.”
- **Charge for it where you can.** A paid pilot filters out tourists and creates internal commitment. Credit it against year one if that helps the conversation.
- **Agree what happens if it succeeds** before it begins. “If we hit those numbers, what’s the process to go ahead?” Ask this at the start, not at the end.

Closing without pressure tactics

You don’t need artificial urgency and you shouldn’t use it — you’re a small company and your credibility is your main asset. What you need instead is a plan that both sides can see.

Write a **mutual action plan**: a short shared document listing every step from here to live, who owns each one, and the date. Security review, legal, the internal approval meeting, the kickoff. Send it after the demo and keep it updated.

This does three things at once. It surfaces steps you didn’t know about — usually procurement and security, which are the two things that quietly add six weeks. It makes slippage visible to your champion rather than only to you. And it turns “just checking in” follow-ups, which are annoying and low-status, into “we’re at step 4, are we still good for the 18th,” which is professional and hard to ignore.

WORKSHEET 6

Mutual action plan

#	STEP	OWNER	DATE
1	Discovery complete, pain quantified		
2	Demo with economic buyer present		
3	Pilot scope + success criteria signed		
4	Security / IT review		
5	Pilot mid-point review		
6	Results reviewed against criteria		
7	Commercials agreed		
8	Legal / procurement		
9	Signature		
10	Kickoff and first value		

Losing well

You will lose most deals. The ones you lose are worth almost as much as the ones you win, if you extract the reason — and “budget” is never the reason, it’s the polite wrapper around the reason.

Two weeks after a loss, send a short note: “No pitch, and I’m not trying to reopen it. I’m trying to get better at this. What was the actual reason we didn’t go ahead?” The response rate is surprisingly high, precisely because you’re not selling. Log every answer in one place and read them monthly. Patterns will appear within about fifteen losses, and those patterns are your product roadmap and your messaging backlog.

The minimum pipeline system

You need a CRM. Not an expensive one — a free tier or a well-structured spreadsheet is fine at ten deals. What matters is discipline, not software.

Six fields, updated after every conversation

- Company, contact, and role
- Stage (the six from Figure 6)
- Expected value and expected close date

- The quantified pain, in their words
- Next step and its date — blank means the deal is dead, treat it as dead
- Where it came from (channel), because Chapter 10 needs this and you will not remember

The weekly review, 30 minutes, alone

Every Friday: any deal with no next step gets closed out or a next step gets created. Any deal whose close date has slipped twice gets a hard conversation or gets marked lost. Count how many new conversations you had and compare it to last week. That's it. The value of this ritual isn't the data, it's that it stops you from carrying dead deals in your head as if they were real.

TWO NUMBERS WORTH KNOWING ABOUT YOUR OWN DEALS

In 2026 benchmarks, average B2B software sales cycles ran about **6 weeks under \$10K ACV, 12 weeks at \$10–50K, 17 weeks at \$50–100K, and 24 weeks above \$100K.**¹⁶ And opportunities sourced by the seller won at about 38%, versus 23% for marketing-sourced — while opportunities sourced from existing customers won at 52%.¹⁷

Two implications. Your deals are probably not stalled, they're just the length they are — check your expectations against the band. And the highest-converting pipeline in the business comes from people who already pay you, which almost nobody at your stage works systematically.

The seven mistakes I see most

1. **Demoing too early.** The demo is the reward for a good discovery call, not a substitute for one.
2. **Selling to the user and never meeting the buyer.** Enthusiasm from someone with no budget feels identical to progress and is not progress.
3. **Discounting to close.** It teaches the customer your price is fiction and it doesn't fix a value problem.
4. **Building a feature to win one deal.** Occasionally right. Usually a distraction dressed as customer focus. Ask whether three other Tier A accounts need it.
5. **Not writing anything down.** If it's in your head, you can't improve it, delegate it, or spot the pattern.
6. **Treating “very interested” as a stage.** Interest is not a commitment. Only dated actions are.
7. **Stopping outreach when deals get busy.** The classic founder sine wave: prospect, get busy, close, panic, prospect. Ring-fence the outreach block even in a good month, especially in a good month.

THE OUTPUT OF THIS CHAPTER

Not revenue — though you should get some. The output is a written document that says: *this kind of company, contacted this way, with this message, agrees to a call about X% of the time, and about Y% of those become customers in about Z weeks, for about £N.*

That paragraph is the thing you can scale, hire against, and raise money on. Twenty conversations is usually enough to write the first draft of it.

Pricing and Packaging

Most early founders set their price by opening three competitor pricing pages, picking something slightly lower, and never touching it again. It's understandable and it's expensive — pricing is the only lever that moves revenue with no delivery cost attached, and it's usually the least-examined decision in the company.

You are almost certainly too cheap. Not because confidence is a strategy, but because founders anchor on what they'd personally pay, and you are not your customer — you're a person who knows how the sausage is made and therefore systematically undervalues it.

Start with the value metric, not the number

The value metric is *what you charge for*. Per seat, per document, per workflow run, per location, per GB, per closed deal. It matters more than the price itself, because it determines whether your revenue grows as your customer gets more value — or whether you get punished for being useful.

The test is one sentence: **can you explain in a single line why the bill goes up as the customer gets more value?** If not, the metric is wrong.

METRIC	WORKS WHEN	BREAKS WHEN
Per seat	Value scales with the number of people using it daily	Value comes from automation — you're charging for headcount you're removing. Very common mismatch in AI products.
Per unit of work (document, run, ticket, transaction)	Value clearly tracks volume, and volume is visible to the buyer	Volume is unpredictable, so buyers can't budget and procurement stalls
Per outcome	The outcome is unambiguous and attributable to you	Attribution is contested, or the customer influences the outcome as much as you do
Flat platform fee	Value is "this exists and works," and usage varies wildly	You have a whale using 40× the average and paying the same
Hybrid (base + usage)	You want predictable revenue and upside. Most common in 2026.	You make it complicated enough that buyers can't self-estimate their bill

The mismatch to watch for in AI products

If your product replaces work that a person used to do, per-seat pricing is actively working against you. The better your product gets, the fewer seats the customer needs, and your revenue falls as your value rises. Charge for the work done, or for the outcome, or use a base plus usage — but don't charge per human for a product whose whole argument is that fewer humans are needed.

What the market is actually doing

Pricing model mix among B2B software companies going into 2026, from a survey of GTM leaders at 150+ companies.¹⁸

FIGURE 8 — PRIMARY PRICING MODEL, 2026 PROJECTED

MODEL	2025	2026P	READ
Hybrid (base + usage)	48%	48%	The default. Predictability plus upside.
Subscription / platform	24%	22%	Slowly declining
Seat-based	14%	15%	Stable, but under pressure from AI-native products
Consumption / usage	9%	9%	Higher among products with proprietary models (~16%)
Outcome-based	5%	6%	Growing, still niche, hard to operationalise

Two things to take from that table. First, hybrid isn't a compromise — it's the consensus answer, and if you're agonising between subscription and usage the market has largely decided you should do both. Second, consumption and outcome pricing are notably more common among products built on proprietary models than among application-layer products,¹⁸ which makes sense: you charge for consumption when consumption costs you money.

A note on committed versus pay-as-you-go

Where companies do use consumption pricing, most attach some upfront commitment rather than pure pay-as-you-go — and pure pay-as-you-go becomes more common as companies get larger and their billing infrastructure matures.¹⁹ The practical translation for you: pay-as-you-go sounds friendly and it needs usage instrumentation, billing plumbing and forecasting discipline you probably don't have yet. Start with a committed minimum plus overage. It's simpler to build, easier to forecast, and buyers under \$100K ACV mostly accept it.

How to price when you have no data

The academic answer is a Van Westendorp price-sensitivity study, which asks buyers four questions about when a price feels too cheap, cheap, expensive and too expensive, and derives an acceptable band.²⁰ It's a good tool and it needs a reasonable sample of real prospects, which you don't have at ten customers. File it for later.

What to do instead, in this order:

1. Build the alternative's cost

Work out what the customer currently spends on the problem — hours × loaded salary, agency fees, error costs, the tool they'd otherwise buy. That number is your ceiling reference. If a broker spends two days per renewal and does forty renewals a month, you can compute what that costs and price as a fraction of it.

2. Take a fraction of the value

Common practice is to price somewhere in the region of 10–25% of the quantified annual value you create. Below 10% and you're leaving money on the table without buying much extra conversion. Above about 30% and the buyer starts doing the maths in a way that doesn't favour you. This is a heuristic, not a law — but it beats guessing.

3. Sanity-check the segment's budget reality

A 15-person agency does not have a \$50K line item, no matter how good the value case is. Price has to fit the buyer's approval threshold as much as their value calculation. Find out what a manager in your ICP can approve without escalating — often something like \$1K/month — and consider deliberately landing under it.

4. Ask, in discovery

You already have a call structure that gets to money. Add these:

"If this worked exactly as we've described, what would that be worth to you in a year?"

"What do you spend on this today — tools, people, time?"

"What would you have to sign off, and who signs it?"

"If I told you it was £2,000 a month, what would your reaction be?" — *then say nothing.*

That last one is uncomfortable and enormously informative. Watch the face. "That's a lot" and "oh, that's fine" are both useful; the second one means you're too cheap.

5. Test it live

Quote a higher number to your next three prospects than you're comfortable with. Not a lie — an actual price you'd honour. Founders discover their price is 2–3× too low far more often than the reverse. If all three say yes immediately, you've learned something worth more than the deals.

THE MOST COMMON EARLY PRICING ERROR

Pricing off your costs. Your infrastructure bill is irrelevant to what your product is worth — it only sets a floor you must clear. A customer paying you £2,000 a month to recover £30,000 a month does not care that it costs you £40 in compute. Cost-plus pricing is how you end up with a business that works and doesn't grow.

Packaging

Packaging is deciding what goes in which box. Three principles cover most of it at your stage.

Three tiers, maximum

Two if you can. Every additional tier multiplies the decisions the buyer has to make and the internal debates they have to have. The classic structure works because it works: an entry tier that's obviously not

enough for a serious buyer, a middle tier where you want everyone, and a top tier that exists partly to make the middle one look reasonable.

Fence on value, not on cost

What separates tiers should be something that correlates with how much the customer gets out of you — volume, users, advanced workflow, integrations, support level. Never gate something that makes the product worse to evaluate: don't put the integration that proves your value behind the top tier, or nobody will ever see the value.

Design for land-and-expand

Your entry point should be small enough to buy without a committee and structured so that success naturally increases spend. This matters more than it used to: retention and expansion are increasingly how software companies are judged, and account executive compensation is visibly shifting from gross new revenue toward net new recurring revenue and net dollar retention.²¹ Even as a founder with no sales team, build the expansion path in now — another team, another location, more volume, a second use case.

The rules of discounting

You need a policy before you need it, because in the moment you will discount.

- **Never discount for nothing.** Every reduction buys something: an annual commitment, a case study, a reference call, a faster close, a bigger scope.
- **Discount the term, not the price.** “20% off if you pay annually” is a cash-flow trade. “20% off because you asked” is a permanent repricing.
- **Hold a floor and say it out loud.** “I can do 15% for annual, and that's the end of what I can do” is more credible than a slow slide, and buyers respect a limit.
- **Watch out for the first big logo.** Discounting 60% to land a name you can put on your site is sometimes right. It is also how companies end up with a flagship customer who costs more to serve than they pay, and a price list nobody believes.

One more, specific to consumption models: if you sell usage-based and the usage doesn't materialise, you have a revenue problem that shows up two quarters later. Companies at scale handle this with commission clawbacks — typically applied three to six months after close for deals above \$50K ACV.²² You don't have commissions yet. But you should have the equivalent discipline in your own forecasting: don't count consumption revenue you haven't seen consumed.

Raising prices

You'll need to, probably within a year, and it's less traumatic than you fear. Three ways to do it without a fight:

1. **New customers only, immediately.** Change the price, grandfather everyone existing. Costs you nothing and tells you within a month whether the market notices.
2. **Existing customers at renewal, with notice and a reason.** “We've added X and Y; from your renewal in March the price is Z.” Most churn from price rises comes from surprise, not from price.
3. **Repackage.** Introduce a new tier structure and migrate people to it. Cleanest route when your value metric was wrong in the first place.

The signal that you’re too cheap: prospects say yes too fast and nobody negotiates. In healthy B2B pricing, some proportion of buyers push back. If literally none do, your price is below the point where anyone has to think about it, which means you’re funding your customers’ margin.

WORKSHEET 7

Pricing decisions

1. VALUE METRIC — WE CHARGE PER:

One sentence on why the bill rises as value rises:

2. THE ALTERNATIVE’S ANNUAL COST TO THE CUSTOMER

3. QUANTIFIED ANNUAL VALUE WE CREATE — AND OUR TARGET SHARE OF IT (10–25%)

4. APPROVAL THRESHOLD OUR CHAMPION CAN SIGN WITHOUT ESCALATING

5. THE LADDER

TIER	PRICE	WHO IT’S FOR	THE FENCE
Entry			
Core			
Scale			

6. DISCOUNT POLICY — MAXIMUM, AND WHAT WE GET FOR IT

7. EXPANSION PATH — HOW AN ACCOUNT NATURALLY DOUBLES

IF YOU DO ONE THING FROM THIS CHAPTER

Quote your next three prospects a price 50% higher than your current one. Keep everything else identical. You will learn more about your pricing in three calls than in three months of reading, and the downside is bounded — you can always come down, and coming down is a concession you can trade for something.

PART THREE

The Engine

You’ve proven you can sell it. Now the question changes from “does this work?” to “how do I make this happen again on purpose, more often, without me in every conversation?”

Four decisions: how people buy from you, where they come from, what the numbers have to be, and what you hand to a machine. Get these right and you have a company. Get them wrong and you have a founder doing sales forever with a slowly growing sense of dread.

-
- | | |
|----|---|
| 08 | Choosing Your Motion — self-serve, sales-led or hybrid, decided by ACV rather than preference |
|----|---|
-
- | | |
|----|--|
| 09 | The Five Doors — where pipeline comes from, and how to test a channel properly |
|----|--|
-
- | | |
|----|--|
| 10 | The Numbers — your GTM model on one page, and the benchmarks to judge it against |
|----|--|
-
- | | |
|----|---|
| 11 | AI as Leverage — what it measurably improves, and the three places it quietly costs you |
|----|---|
-

Choosing Your Motion

“Motion” means how a customer gets from never having heard of you to paying you. There are three, they cost wildly different amounts to run, and the right one for you is mostly determined by a single number: how much a customer is worth in year one.

This is not a taste decision. Founders pick sales-led because they like control, or self-serve because they dislike calls. Both of those are how you end up with a motion your economics can't support.

The three motions

MOTION	HOW IT WORKS	ECONOMICS
Self-serve (product-led)	Someone finds you, signs up, gets value alone, and pays with a card. No human involved unless something breaks.	Low cost per customer, needs volume. Fast payback but typically weaker retention.
Sales-led	A human runs discovery, demo, pilot, negotiation. Buyer never touches the product unsupervised until they're a customer.	High cost per customer, works at higher ACV. Longer payback, stronger retention.
Hybrid	Self-serve entry, human involvement at the point where deal size or complexity justifies it.	The 2026 default. Harder to run, better economics than either alone.

Let ACV decide

Here's the blunt version. If a customer is worth \$600 a year, you cannot afford a human conversation with them — one hour of your time already exceeds the gross margin on year one. If a customer is worth \$60,000 a year, expecting them to sign up unassisted and expense it is fantasy; someone in procurement is going to want a call.

Sales cycle length scales with deal size in a very consistent way. Across 150+ B2B software companies in 2026, average cycles ran roughly:¹⁶

6 weeks — under \$10K ACV	12 weeks — \$10K–\$50K	17 weeks — \$50K–\$100K	24 weeks — \$100K+
-------------------------------------	----------------------------------	-----------------------------------	------------------------------

Combine that with cost per opportunity — averaging around \$6,300 for SMB, \$8,000 for mid-market and \$11,200 for strategic accounts in 2026²³ — and the decision mostly makes itself.

FIGURE 9 — MOTION BY ANNUAL CONTRACT VALUE

ACV	DEFAULT MOTION	WHAT IT LOOKS LIKE IN PRACTICE
Under \$1K	Pure self-serve	No calls, ever. Everything is product, docs, and email. Your marketing has to do 100% of the convincing. Needs real volume to matter.
\$1K–\$10K	Self-serve + assist	Sign up alone, but a human appears for the accounts that look valuable. Sales is a conversion aid, not a gate.
\$10K–\$50K	Hybrid, sales-forward	Free trial or POC as the qualifier, then a real sales process. This is where most B2B startups actually live.
\$50K+	Sales-led	Multi-stakeholder, security review, procurement, pilot with success criteria. You personally, for a long time.

The affordability check

Rough gate for adding a human to a deal:

Year-1 gross profit per customer > 3 × cost of the sales time

e.g. \$12,000 ACV × 80% margin = \$9,600 gross profit

6 hours of selling at a fully loaded \$150/hr = \$900

9,600 / 900 = 10.7× → comfortably affordable

e.g. \$900 ACV × 80% = \$720 gross profit

3 hours of selling = \$450

720 / 450 = 1.6× → you are losing money on the motion

At your stage your own time feels free. It isn't — it's the scarcest input in the company. Run the check with a real hourly number.

Hybrid isn't a hedge anymore

A few years ago “hybrid” meant you hadn't decided. In 2026 it's the mainstream answer. Around 60% of B2B software companies describe their customer acquisition as hybrid rather than purely top-down or bottom-up, and high-growth companies lean into bottom-up entry points harder than their peers — expecting roughly 19% of revenue from self-serve versus about 8% for everyone else.²⁴

The pattern that's working: **self-serve as a wedge, sales as the expansion**. An individual user adopts the product on their own credit card or a free tier. Usage spreads inside the company. That footprint becomes the evidence for a top-down conversation with an executive about a broader deal. You're not selling into a cold account — you're selling into an account where forty people already use you.

Cursor is the extreme version of this. It grew from roughly \$4M to \$100M ARR in about twelve months on organic developer adoption with no paid marketing and no outbound sales, then expanded inside companies like NVIDIA and Salesforce bottom-up, and only brought in a senior sales leader after passing \$2B ARR — at which point sales formalised demand that already existed rather than creating it.²⁵

You are not Cursor. Nobody is. But the sequence is instructive and it generalises: *earn adoption first, add sales to capture the value of adoption you already have*. Adding sales to a product nobody adopts organically doesn't produce Cursor's outcome; it produces expensive silence.

If you go self-serve: the four front doors

Free trial, freemium, reverse trial, or interactive demo. The differences matter more than people assume, and the conversion data is counterintuitive in one important way.

MODEL	ADOPTION	FREE-TO-PAID	USE IT WHEN
Free trial no card	57% of products use a trial ²⁶	~5% — low, but high signup volume	Value is obvious within days and you want maximum top-of-funnel
Free trial card required	—	~30% — roughly 6x higher ²⁶	You'd rather have fewer, more serious trials. Usually the right call for small teams
Freemium	26%	2–8% typical	The free tier is genuinely useful, and usage creates network or data effects
Reverse trial	7%	—	Everyone starts on the paid tier, drops to free when it expires. Good compromise.
Interactive demo	7%	—	Product is complex to set up, or the buyer can't get data into it quickly

Median free-to-paid conversion across around 200 B2B products surveyed in January 2026 was about 8%.²⁶ The counterintuitive part: free trials convert at higher rates than freemium, but the difference largely disappears once you account for how many more people sign up for a freemium product in the first place. Conversion rate is not the metric. **Paying customers per thousand visitors is the metric.**

REQUIRING A CARD IS THE MOST UNDER-USED LEVER IN EARLY SELF-SERVE

It cuts your signup count and roughly multiplies your conversion rate, which for a small team is almost always the better trade — you get fewer, better users, less support load, and a much clearer read on whether the product works. You can always remove the requirement later when you need volume. Adding it later feels like a downgrade to your existing users.

Activation is the whole game

If you take one thing from this chapter: whether someone becomes a paying customer is decided almost entirely by whether they reached the moment where the product visibly worked for them. Not by your pricing page, not by your email sequence.

The gap is enormous. Activated trials convert to paid at rates in the tens of percent; un-activated trials convert in low single digits.²⁷ Everything else you might optimise is rounding error next to that.

Defining your activation moment

It's the first time the user sees the product do the valuable thing with their own stuff. Not signup. Not completing a profile. Something like: connected their inbox and saw three real drafts. Uploaded a document and got the extracted data back. Ran the workflow once and it worked.

Write yours down as a single measurable event, then instrument it. Then find out two things: what percentage of signups reach it, and how long it takes them. Those two numbers are the most valuable telemetry in an early self-serve product.

Getting more people there

- **Cut steps.** Every field, confirmation and configuration screen between signup and value costs you users. Ask what you'd delete if activation rate were the only thing you were paid on.
- **Bring the value forward.** Show something useful before you ask for setup. Pre-populate with sample data that resembles theirs.
- **Do it for them.** At low volume, manual onboarding is a superpower, not a failure. Get on a call and set it up yourself. You'll learn where the product loses people, and every one of those observations is a permanent fix.
- **Instrument the drop-off, then watch five recordings.** Not a dashboard — actual sessions. Painful and clarifying.

And note that self-serve doesn't mean human-free. Around 80% of trial-based B2B products include human touchpoints for their larger prospects.²⁶ The product handles volume; you handle the accounts worth handling.

WORKSHEET 8

Motion selection

Our target year-1 ACV

Implied motion (Figure 9)

Gross margin per customer, year 1

Hours of human time per closed deal

Affordability ratio (target > 3×)

Front door: trial / freemium / reverse / demo

Card required at signup? Y / N — and why

Activation event — one measurable moment

Current % of signups reaching it

Current median time to reach it

Expansion trigger — when a human gets involved

If you can't fill in the activation event, stop and do that before anything in Chapter 9. Pointing channels at a product people can't get started with is how marketing budgets disappear.

The Five Doors

Pipeline comes from five places. Not fifty. Once you see them as five doors rather than an endless list of tactics, channel strategy stops being overwhelming and becomes a small set of bets you can size, test and kill.

Before the doors, one number that should reframe how you think about this. Among high-growth B2B software companies under \$100M revenue, roughly **62% of new-logo pipeline comes from sales-sourced motion and 19% from channel and partnerships** — against about 12% from marketing. Their slower-growing peers get 47% from sales and 34% from marketing.²⁸

62% of new-logo pipeline from sales-sourced motion at high-growth companies under \$100M	19% from channel and partnerships	12% from marketing — versus 34% at slower-growing peers
---	--	--

Read that carefully, because it's easy to draw the wrong conclusion. It does not mean marketing doesn't matter — marketing is doing the work that makes the seller-led motion land, and pipeline attribution systematically undercounts it. What it means is that at your size, **somebody deliberately going and finding customers outperforms waiting for customers to find you**. The fastest-growing small companies are not sitting behind a content calendar.

Door 1 — Outbound

Deliberately contacting people who fit your ICP. It has a bad reputation, largely earned by a decade of people blasting ten thousand identical emails. Done the modern way it's the most controllable pipeline source you have: you decide who, when, and what's said, and you get a signal within two weeks.

Signal-based, not volume-based

The difference between outbound that works and outbound that doesn't is whether your message references something specifically true about that company right now. Generic cold email reply rates now average in the low single digits, while outbound anchored on a real trigger — a funding round, a relevant hire, a technology change, a public complaint — reports reply rates several times higher and a materially lower cost per qualified meeting.⁹

The triggers are the ones from Chapter 3. The discipline is: *no trigger, no send*. Twenty-five researched emails to accounts where something just happened will beat a thousand sends to a bought list, and they won't destroy your domain in the process.

Infrastructure, before you send anything

Get this wrong and nothing else matters, because your mail lands in spam and you won't know for six weeks. The current baseline:

- **Never send cold mail from your primary domain.** Buy lookalike sending domains and forward them to your main site.
- **Authenticate properly.** SPF, DKIM and DMARC set up on every sending domain, no exceptions. Mailbox providers now enforce this for bulk senders.
- **Warm up for at least 30 days** before real volume.
- **Cap around 30 emails per inbox per day.**⁹ Add inboxes for volume; don't push a single one harder.
- **Watch your spam complaint rate** and keep it well under 0.3%. Bounces above 2–3% mean your list is bad — stop and verify.
- **Plain text, no images, no tracking pixels, one link at most.** Every ornament costs deliverability, and a founder's real email doesn't have a header banner.

What to actually write

Short. Specific. About them. The structure that works is four lines:

Subject: something factual and lowercase. "renewal prep at [Company]"

Hi [Name] — saw you're hiring a second claims coordinator.

Most brokerages your size add that role because renewal prep eats two days per account and it's all re-keying carrier PDFs. We read those PDFs and write straight into Epic, so three brokerages we work with cut that to under an hour.

Worth a look, or is this already handled?

Four lines. The trigger, the insight, the proof, an easy out. "Or is this already handled?" gets more replies than any call-to-action, because it's easy to answer and it doesn't presume.

Rules that hold up: no paragraph over three lines, no "I hope this finds you well," no calendar link in the first email, no "quick question" subject line, and never ask for thirty minutes when you can ask a question that costs them nine seconds to answer.

Sequence and channels

Four to six touches over about three weeks, mixed across email and LinkedIn, each one adding something rather than repeating. Touch one: the trigger email. Touch two: a LinkedIn connection with no pitch. Touch three: something genuinely useful — a relevant number, a teardown of their public process, an article. Touch four: a short bump. Touch five: the polite close — "I'll stop here, happy to circle back if this changes." That last one reliably produces replies, including a decent number of yeses.

REALISTIC EXPECTATIONS FOR 100 RESEARCHED ACCOUNTS

On a well-targeted, trigger-led campaign you might see something like 15–25 replies, 5–10 conversations and 1–3 opportunities. Those are illustrative but the shape is right: outbound is a funnel with a wide top and a narrow bottom, and the way to improve it is almost always better targeting, not better copy. When a campaign fails, look at the list before you look at the words.

Door 2 — Inbound

People finding you and raising their hand. Slower to start, compounds forever. The critical distinction, which Chris Walker has spent years hammering on, is between **demand capture** and **demand creation**.²⁹

- **Capture** is meeting people who are already looking: search, comparison pages, review sites, category keywords, branded search. High intent, small audience, competitive.
- **Creation** is making people aware they have a problem worth solving before they start looking. Content with a point of view, talks, podcasts, teardowns, original data.

Remember the 95:5 rule.³ Capture reaches the 5% who are in-market now. Creation is how you get into the heads of the 95% who will be in-market later — and most of that influence is untrackable, which is why it gets underfunded by anyone who manages purely to attribution dashboards.

Do both. But if you have to sequence: build capture for the small number of high-intent queries that exist in your niche — comparison pages, “alternative to X,” the specific problem phrase — and spend the rest of your energy on creation, because that’s what makes the outbound in Door 1 work.

Search is now two systems

This changed recently and most companies haven’t caught up. Buyers increasingly begin with an AI assistant rather than a search box, and generative AI now sits at the top of buyer research interactions in Forrester’s 2026 data.⁴ When someone asks an assistant for the best tools for their situation, the handful of names it returns become the shortlist, and everyone else is invisible.

Ranking on Google does not automatically get you there. In one analysis of B2B SaaS buyer queries, only about 12% of pages ranking in Google’s top three were actually cited by the major assistants; the ones that missed out lacked clear answer blocks, unambiguous entity definitions, and third-party validation.⁵

WHAT ACTUALLY HELPS YOU GET CITED

- **Answer the question in the first two sentences**, then explain. Assistants extract answers, not narratives.
- **Be unambiguous about what you are and who you serve** on your own site — in plain sentences, not a hero animation. Models learn your entity definition from your own copy and from how others describe you.
- **Get described accurately in third-party places** — review sites, directories, comparison articles, podcasts, forum threads. Independent corroboration is what models weight.
- **Publish comparison and alternative pages** honestly, including who you’re not for. These get cited disproportionately because they answer the exact question buyers ask.
- **Structure ruthlessly**. Headings that match real questions, short paragraphs, lists, tables, dates.
- **Check what the assistants currently say about you**. Ask four of them, in a fresh session, the questions your buyer would ask. If they describe you wrongly, that’s a live problem — most buyers take the answer at face value and never click through.

What to publish if you’re technical and time-poor

Not a blog with weekly SEO posts about industry trends. That’s a treadmill built for teams with a content budget. The formats with the best return for a founder are:

- **Original data from your own product or research.** Nobody else has it, it gets cited and linked, and it doubles as sales proof.
- **Teardowns.** Take a real, specific situation in your niche and analyse it publicly. This is the fastest way to demonstrate expertise without asking anyone to trust you.
- **The definitive page on one narrow problem** — the one your ICP searches at 11pm. One excellent page beats forty mediocre ones for both search and citation.
- **Free tools and calculators.** High effort, high durability, and they capture the exact audience with the exact problem.
- **Documentation.** Underrated as marketing. Technical buyers evaluate you through your docs before they ever talk to you.

Door 3 — Founder-led distribution

You, in public, consistently. It's the cheapest asymmetric advantage a small company has, because personal accounts consistently outreach company accounts and because people trust a person more than a logo.

What works, concretely: three to five substantive posts a week on the platform where your buyers are — usually LinkedIn for B2B, sometimes X, occasionally a specific forum. Quality over cadence; the algorithm rewards substance and punishes filler more than most people assume.³⁰

Three formats that reliably earn attention:

- **Building in public.** What you shipped, what broke, what you learned, what a number actually was. Specific, honest, unglamorous.
- **The contrarian take.** Your point of view from Chapter 4, applied to something current. Genuine disagreement, not manufactured hot takes.
- **Meeting aftermath.** “Spoke to four ops leads this week. Three of them said the same thing about X.” Real signal from real conversations, anonymised.

Two warnings. First, this takes months, not weeks — the compounding is real but slow, and most founders quit at week six. Second, it has a ceiling: founder-led social tends to stop being sufficient as the primary motion somewhere around \$5–10M ARR, at which point it has to become one voice among several rather than the whole engine.³¹ That's a good problem and it's a long way off.

Communities deserve a specific mention because founders get them wrong. Join the three places where your ICP actually talks. Then read for a month and answer questions with no mention of your product. The return on being genuinely useful in a niche community is enormous and it is destroyed permanently the first time you pitch.

Door 4 — Partnerships and channel

Consistently the most under-used door at early stage, and the data suggests that's a mistake — high-growth companies under \$100M pull about 19% of new-logo pipeline through channel and partner motion, meaningfully more than their slower peers.²⁸

Four kinds, in order of how quickly they pay off for a small team:

1. **Integration partners.** The platforms your ICP already lives in. Being in their marketplace or directory puts you in front of a pre-qualified audience at the moment they're looking for exactly your category.
2. **Service partners.** Agencies, consultancies, bookkeepers, implementation firms — whoever your customers already pay for advice. They have trust you'd take years to build and an incentive to bring you in if you make them look good.
3. **Complementary vendors.** Non-competing products selling to the identical buyer. Co-hosted sessions, shared content, warm introductions both ways.
4. **Resellers.** Real leverage, real overhead. Usually premature under \$2M ARR.

How to start, practically: pick *one* partner and do co-selling, not co-marketing. Co-marketing means a joint webinar that produces a list of people who wanted the recording. Co-selling means you sit on a call with their client and solve something. Do that five times with one partner before you build a partner programme, a portal, or a deck.

Door 5 — The customers you already have

The highest-converting pipeline in B2B, and almost nobody at your stage works it deliberately. Opportunities sourced through customer success won at about 52% in 2026 — against 38% for sales-sourced and 23% for marketing-sourced.¹⁷ Of course they do: you're selling into a relationship where value is already demonstrated.

Three things to run on a schedule:

- **The referral ask, specifically.** Not “know anyone who'd like this.” Instead: “You mentioned you came from [Company]. Is there anyone there who has the same problem you had?” Named, specific, easy to say yes to. Ask at the moment of visible success, not at renewal.
- **Expansion.** Another team, another site, another use case, more volume. Look at your usage data for accounts pressing against a limit and go and talk to them — that's a warm opportunity you're currently letting expire.
- **The quarterly value conversation.** Thirty minutes with each customer: what's better, what's still broken, what's next. It reduces churn, produces case studies, and generates the referral and expansion conversations above without needing a separate motion.

And a note on why this matters more than it used to: contract lengths are shortening, with sub-one-year deals tripling as a share of new-logo subscriptions between 2025 and 2026.¹² Shorter contracts mean retention is re-earned constantly. The install base isn't a passive asset any more — it's a pipeline source and a risk, both of which need someone's attention. Right now that's you.

Where paid advertising fits

Briefly, because founders overrate it. Paid works well for two things at your stage: capturing existing high-intent demand (searches for your category, your competitors, the specific problem) and staying in front of people who already know you (retargeting site visitors and your target account list).

Paid works badly for creating demand from cold audiences when you have no brand, no case studies and a \$2,000 test budget. You will spend it, learn very little, and conclude ads don't work. They do — at a spend level and a stage that isn't yours yet.

Choosing: pick three, test, focus on one

You cannot run five doors. You probably can't run three properly. The Bullseye method from *Traction* is the cleanest way to decide: brainstorm how each channel could work for you, sort them into promising / possible / unlikely, pick three from the promising pile, run cheap parallel tests, then put everything into the one that shows real traction.³²

The failure mode is skipping the test and going straight to commitment based on what you'd enjoy doing. The other failure mode is testing forever. A test has an end date.

WORKSHEET 9

Channel test card — one per channel, filled in before you start

Channel

Hypothesis — "If we do X for Y, we expect Z"

Minimum volume for a valid signal

Budget (money)

Budget (hours per week)

Run time — start and end date

Leading metric we'll watch weekly

Success looks like

Kill criterion — what makes us stop

What we'll learn even if it fails

Rules of thumb: minimum four weeks for outbound, eight for paid, six months for content and founder-led social. If you can't commit to the minimum, don't start the test — you'll get a false negative and cross off a channel that would have worked.

THE REALISTIC ANSWER FOR MOST READERS OF THIS BOOK

Two doors. **Outbound**, because it's controllable, fast to read, and it's what the fastest-growing small companies actually over-index on. And **founder-led distribution**, because it compounds, costs only your time, and makes the outbound work by the time it reaches someone who has already seen your name.

Add partnerships as door three once one of the first two is producing reliably. Everything else waits.

The Numbers

At some point “we need more customers” has to become “we need 340 more of the right conversations this quarter.” That translation is what this chapter does. It’s arithmetic, it takes an afternoon, and it changes what you do on Monday more than any other single exercise in this book.

The purpose isn’t forecasting accuracy — your forecast will be wrong. The purpose is to find out which part of your funnel is actually broken, and to stop optimising the part that isn’t.

The shape: it doesn’t end at the sale

The old mental model is a funnel: wide at the top, narrow at the bottom, done when money changes hands. That model is wrong for subscription software, and Winning by Design’s bowtie is the correction — the funnel mirrors after the close, with onboarding, adoption, renewal and expansion as an equally instrumented second half.³³

FIGURE 10 — THE BOWTIE

HALF	STAGES	THE QUESTION
Acquisition	Reach → Engaged → Qualified → Closed-won	How efficiently do strangers become customers?
Retention	Onboarded → Impact delivered → Renewed → Expanded	How much of that customer's lifetime value do we actually collect?

Why this matters commercially rather than conceptually: the right half is where most of the money is, and it’s the half founders instrument last. Top-quartile companies under \$50M ARR ran net dollar retention of around 121-123% through 2025³⁴ — meaning the existing customer base alone grew revenue more than 20% a year with no new logos at all. If your NDR is under 100%, you are refilling a leaking bucket and every pound you spend on acquisition is worth less than it looks.

Working backwards from the number

Pick a target. Say you want to add \$300,000 of new ARR in the next twelve months, and your average contract value is \$12,000.

Target new ARR	\$300,000
÷ ACV	\$12,000
= Customers needed	25
÷ Win rate (opportunity → won)	25%
= Qualified opportunities needed	100
÷ Meeting → opportunity rate	40%
= Meetings needed	250
÷ Contact → meeting rate	6%
= Contacts needed	4,167
Per month:	~350 contacts, ~21 meetings, ~8 opportunities, ~2 customers
Per week:	~80 contacts, ~5 meetings, ~2 opportunities

Now the model has told you something specific and slightly alarming: you need about eighty properly researched contacts a week. That's a real workload, it's schedulable, and it converts a vague goal into a calendar.

It also shows you the leverage points. Improving your win rate from 25% to 35% removes about 1,200 contacts from the year. Raising ACV from \$12,000 to \$18,000 removes eight customers from the target. In almost every early-stage model, **ACV and win rate are cheaper to improve than volume** — and both of them come from Part I of this book, not from working harder.

USE YOUR OWN RATES AS SOON AS YOU HAVE TWENTY CONVERSATIONS

Start with benchmark rates because you have nothing else. Replace each one with your actual number as soon as you've seen it happen twenty times. Your rates will differ from the benchmarks, sometimes by a lot, and the differences are the most interesting thing in your business — they tell you where you're unusually good and where you're quietly bleeding.

Pipeline coverage

Coverage is the ratio of open pipeline value to the number you need to close. If you need \$100,000 this quarter and have \$300,000 in live opportunities, you're at 3×

The mathematically correct coverage is $1 \div \text{your win rate}$. A 25% win rate needs 4× to be certain. In practice, teams target around 3× for smaller deals and more as deals get bigger and slower, because larger deals slip:

ACV BAND	COVERAGE	WHY
Under \$25K	3×	Short cycles, less slippage, higher win rates
\$25K–\$100K	3.5–4×	More stakeholders, more slip
\$100K–\$250K	5×	Procurement, security, budget cycles

Common industry practice.³⁵ For reference, companies in the 2026 ICONIQ dataset reported average pipeline coverage of around 3.6–3.9×.³⁶

The important discipline isn't the ratio, it's the honesty. Coverage calculated on a pipeline full of deals with no next step is a number that will lie to you for two quarters and then present you with a bill.

The eight numbers to track before \$1M ARR

Not thirty. Eight. Everything else is interesting and none of it changes a decision at your size.

FIGURE 11 — THE PRE-\$1M DASHBOARD

#	METRIC	WHY IT EARNS A SLOT
1	New conversations per week	The only input entirely under your control. If it drops, everything drops eleven weeks later.
2	Qualified opportunities created	The earliest reliable predictor of revenue.
3	Win rate	Tells you whether the problem is volume or the offer. Most founders assume volume; it's usually the offer.
4	Average contract value	The cheapest lever in the model, and the one nobody pulls.
5	Sales cycle length	Determines how long you have to wait to know anything. Compare to your ACV band.
6	Activation rate	% of signups reaching first value. Governs everything in a self-serve motion.
7	Logo and revenue churn	The number that decides whether you're building an asset or a treadmill.
8	Pipeline by source	Which door produced it. Without this you cannot decide where to spend next quarter.

Two ratios for once you're spending money

CAC payback — how many months of gross profit it takes to recover what you spent acquiring a customer. It's the most useful efficiency measure at early stage because it's about cash, and cash is what you're short of.

$$\text{CAC payback (months)} = \frac{\text{Sales \& marketing spend in period}}{(\text{New ARR added} \times \text{gross margin})} \times 12$$

Commonly cited 2026 medians land somewhere around 15–18 months, with efficient companies under 12 and self-serve motions frequently in the 6–12 month range.³⁷ Treat those as orientation rather than targets — the number that matters is whether yours is getting shorter.

LTV:CAC gets quoted constantly and is close to useless before you have two years of retention data, because LTV is a projection built on a churn rate you've barely measured. Ignore it until it means

something.

Benchmarks, in one place

Use these to sanity-check, not to target. Every one is an average across companies far larger than yours.

FUNNEL CONVERSION, 2026 AVERAGE ACROSS 150+ B2B SOFTWARE COMPANIES¹⁵

STAGE	2025	2026	CHANGE
New lead → MQL	25%	28%	+3
MQL → SQL	27%	30%	+3
SQL → closed-won	25%	28%	+3
Demo → closed-won	36%	38%	+2
Free trial / POC → paid	36%	50%	+14

WIN RATE BY OPPORTUNITY SOURCE, 2026¹⁷

SOURCE	WIN RATE	IMPLICATION FOR YOU
Customer success	52%	Work your install base. Highest-yield pipeline you own.
Sales	38%	Seller-sourced beats marketing-sourced. Do outbound.
Channel / partner	35%	Under-used at early stage relative to its return.
Marketing	23%	Volume source, lower intent. Doesn't mean it's not working.

ACQUISITION COST BY SEGMENT, 2026 AVERAGE²³

SEGMENT	COST PER LEAD	COST PER OPPORTUNITY
SMB	\$500	\$6,300
Mid-market / enterprise	\$600	\$8,000
Strategic	\$800	\$11,200

That cost-per-opportunity column is worth sitting with. Six to eleven thousand dollars to create *one qualified opportunity*, which then converts at roughly 28%. Which means a closed customer costs somewhere in the region of \$22,000–\$40,000 to acquire through a conventional sales-and-marketing motion at those companies.

You almost certainly can't spend that. What you have instead is time, a founder's credibility, and the ability to do things that don't scale — which is precisely why founder-led sales and founder-led distribution are the right answer at your stage rather than a compromise.

The weekly ritual

Thirty minutes, same time every week, one page. Not a dashboard project — a spreadsheet you fill in by hand, because filling it in by hand is what makes you look at it.

FIGURE 12 — THE WEEKLY ONE-PAGER

THIS WEEK	ACTUAL	TARGET	TREND
New conversations started			
Meetings held			
Opportunities created (value)			
Opportunities closed-won (value)			
Opportunities lost — and the reason			
Total open pipeline / coverage			
Signups → activated			
Churn / at-risk accounts			
Deals with no next step (should be 0)			

The last row is the one that keeps you honest. Any deal without a dated next step gets one this week or gets closed out. It's a small ritual and it's the difference between a pipeline and a list of people you once spoke to.

What not to measure

A short list, because measuring the wrong thing is worse than measuring nothing — it produces confident bad decisions.

- **Website traffic** as a headline number. Traffic from people who will never buy is a cost, not an asset. Track traffic from your ICP if you can segment it; ignore the total.
- **Followers and impressions.** Useful as a directional health check on Door 3, meaningless as a goal.
- **Emails sent.** An activity metric that rewards exactly the behaviour that destroys your deliverability.
- **MQs** as a target you hold anyone to. The moment a number becomes a target, the definition erodes to meet it.
- **Last-touch attribution** as truth. It systematically over-credits the final click and erases everything that made the buyer aware of you — which, given the 95:5 rule, is most of the work. Ask people how they heard about you in a free-text field and read the answers. It's cruder and more accurate.

WORKSHEET 10

Your GTM model, one page

INPUT	YOURS	BENCHMARK
Target new ARR (12 months)		—
Average contract value		—
Customers needed		= ARR ÷ ACV
Win rate (opportunity → won)		28–38%
Opportunities needed		
Meeting → opportunity rate		~40%
Meetings needed		
Contact → meeting rate		3–10%
Contacts needed — year / month / week		
Sales cycle (weeks)		6 / 12 / 17 / 24 by ACV
Required pipeline coverage		3–6× by ACV
Gross margin		>70%
CAC payback (months)		<18
Net revenue retention		>100%, top quartile 120%+

The question this page answers: is the weekly contact number a workload a human being can actually do? If it isn't, you don't have a volume problem — you have an ACV problem or a win-rate problem, and both of those get fixed in Part I.

AI as Leverage

You built your product with AI doing a lot of the typing. The obvious next thought is to do the same to your go-to-market. It works — but not evenly, and the data on where it works is more specific and more useful than the marketing around it.

What the numbers actually show

Comparing companies with AI fully embedded in their GTM processes against those where it isn't, across the 2026 ICONIQ dataset:³⁸

FIGURE 13 — HIGH AI ADOPTERS VS EVERYONE ELSE			
MEASURE	HIGH ADOPTION	MED / LOW	READ
Ramped reps hitting quota	67%	59%	Meaningful
SMB quota attainment	106%	80%	Largest gap
Enterprise quota attainment	96%	77%	Large
Cost per opportunity	\$8.5K	\$8.9K	Slightly better
Cost per lead	\$670	\$600	Worse
Sales cycle (weeks)	19.3	19.6	Barely moved
Pipeline coverage	3.9×	3.6×	Better
Late renewals	27%	29%	Slightly better

And on funnel conversion, comparing companies where more than half of pipeline is AI-influenced against those where it's less:³⁹

+11pt New lead → MQL 38% vs 27%	+8pt MQL → SQL 37% vs 29%	+1pt SQL → closed-won 29% vs 28%	+3pt Demo → closed-won 40% vs 37%
--	--	---	--

That pattern is the entire chapter in one picture. AI substantially improves the top of the funnel and changes almost nothing once a deal is live. More opportunities, better qualified, reaching a human faster.

But the conversation where someone decides to spend money is still a conversation, and it converts at roughly the same rate as it always did.

The cost-per-lead line is the honest detail most write-ups skip. High adopters pay slightly *more* per lead while paying less per opportunity — they’re generating more, and better, not cheaper. Volume was never the constraint.

The leverage that actually matters to you

Look at headcount. Total GTM staff — sales, post-sales, marketing and revenue operations — by AI adoption level:⁴⁰

REVENUE BAND	HIGH AI ADOPTION	MED / LOW ADOPTION	DIFFERENCE
\$10M–\$25M	20	35	–43%
\$25M–\$100M	45	65	–31%
\$100M–\$250M	125	165	–24%
\$250M–\$500M	275	350	–21%

A twenty-person go-to-market team doing what used to take thirty-five. Note the gap is largest at the smallest revenue band — which is the direction of travel toward you. For a founder with no team, the practical version of this is: **the work of a small GTM team is now within reach of one determined person with good systems.** Not all of it. The researching, listing, drafting, summarising, logging and monitoring — yes.

What to hand over, and what never to

FIGURE 14 — THE AUTOMATION LINE

AUTOMATE AGGRESSIVELY	WHY IT'S SAFE
Account research and enrichment	Finding facts. High volume, low judgment, easily verified.
List building and scoring	Applying your rubric consistently. Better than you at 4pm on Friday.
Trigger monitoring	Watching job boards, funding news, hiring pages, tech changes. Tedious and continuous.
First drafts of everything	Emails, posts, docs, proposals. You edit; it types.
Call transcription and summary	Frees you to actually listen. Then mine the transcripts for language.
CRM hygiene and logging	The admin that otherwise doesn't get done.
Usage-signal alerts	Telling you an account is approaching a limit or has gone quiet.
NEVER AUTOMATE	WHY IT COSTS YOU
Discovery conversations	The whole point is learning what you didn't know to ask. A model can't be surprised on your behalf.
Deciding your positioning	Requires judgment about a market you're inside and it isn't.
The final word of outbound to Tier A	Your 25 best accounts deserve a human sentence. The machine writes the other 75.
Loss reasons and churn conversations	The most valuable information in your business arrives in these, unstructured and defensive.
Anything requiring a promise	Roadmap commitments, pricing exceptions, guarantees. Never delegate a promise you'd have to honour.

The stack for a team of one to five

Deliberately minimal. The most common failure I see is a founder with \$400/month of tools and eleven customers.

1. **A CRM.** Free tier, or a spreadsheet with the six fields from Chapter 6. Non-negotiable, because your memory is not a system.
2. **A prospecting and enrichment layer** to find companies and contacts, and to run your ICP rubric across a list. This is where most of the AI leverage lives at your stage.
3. **A sending tool with its own domains and warmup.** Separate from your real email.
4. **Call recording with transcripts.** Cheap, and it turns every call into a searchable corpus of your customers' exact language.
5. **Product analytics** good enough to answer: who signed up, did they activate, are they still here.

6. An LLM you actually use properly — with your positioning, ICP, messaging map and call transcripts in context, so its drafts sound like your business rather than like software.

That's it. Add nothing else until a specific bottleneck demands it, and then add one thing.

THE GTM ENGINEER

A role that barely existed three years ago and is now a standard early hire: someone who builds go-to-market *systems* — enrichment pipelines, scoring logic, signal monitoring, automated sequences, CRM plumbing — rather than working accounts by hand. Fast-growing companies hire for it earlier than headcount would suggest, because the leverage arrives before the org chart does.

If you're technical, you are already most of the way to being this person for your own company. That's a genuine and slightly underrated advantage: the parts of go-to-market that used to require a team now require a systems-builder, and you are one.

Three places this quietly goes wrong

1. Everyone sounds the same

When every company generates outreach the same way from the same data, the output converges. Your prospect now receives eleven personalised emails a week that all reference their funding round in the same slightly-too-smooth register. Personalisation that anyone can produce isn't personalisation — it's a new baseline, and the differentiator moves to something a model can't manufacture: a genuinely specific observation, a real opinion, an actual piece of work you did before asking for anything.

The practical rule: use AI to decide *who* and *when*. Be much more careful about letting it decide *what*.

2. Volume kills your deliverability

Automation makes it trivially easy to send far more than your infrastructure can survive. Domain reputation is slow to build and fast to destroy, and once it's gone you're not in the inbox for months. Every constraint in Chapter 9 — the warmup, the per-inbox limit, the bounce rate — exists because someone else already learned this expensively.

3. You stop learning

This is the one that actually matters pre-PMF. If a system researches your accounts, writes your emails, summarises your calls and updates your CRM, you can run a whole quarter of go-to-market without ever forming a first-hand model of your buyer. You'll have activity, data and no intuition — and intuition is the thing that lets you make the non-obvious call about positioning or pricing that changes the trajectory.

Before product-market fit, do the research manually enough times to know what the automation is skipping. Automate after you understand, not instead of understanding.

WORKSHEET 11

Automation audit

List every recurring GTM task. Estimate hours per week. Mark A (automate now), P (partially — machine drafts, you approve), or N (never).

TASK	HRS/WK	A/P/N	TOOL
------	--------	-------	------

Finding and researching target accounts			
---	--	--	--

Writing outbound messages			
---------------------------	--	--	--

Following up			
--------------	--	--	--

Call notes and CRM updates			
----------------------------	--	--	--

Proposals and follow-up documents			
-----------------------------------	--	--	--

Content drafting			
------------------	--	--	--

Monitoring product usage signals			
----------------------------------	--	--	--

Reporting and the weekly one-pager			
------------------------------------	--	--	--

Target: reclaim at least eight hours a week, and spend every one of them on customer conversations. That is the only legitimate reason to automate anything at this stage.

THE SUMMARY

AI gets you more, better-qualified opportunities and lets a very small team run a motion that used to need a department. It does not close deals, decide your positioning, or tell you why you lost. Point it at the top of the funnel and at the admin. Keep the conversations, the judgment and the learning for yourself — for now, they’re still the parts that decide whether the company works.

PART FOUR

Sequence

Knowing what to do and knowing what to do on Monday are different problems. This part turns the previous eleven chapters into an order of operations, and then tells you what happens when the thing you've built starts working and you become the bottleneck.

-
- 12 The First 90 Days — a week-by-week plan that installs the whole engine
 - 13 Scaling Past Yourself — when to hire, who first, and what breaks at each stage
-

The First 90 Days

Ninety days is enough to install a working go-to-market engine from nothing. Not to hit a revenue target — that depends on your cycle length and your market — but to have positioning, an ICP, a message, a channel producing conversations, and a model that tells you what's broken. If you follow this and do nothing else, you'll be further ahead than most companies at your stage.

Assume roughly 15–20 hours a week on go-to-market. If you have more, the plan compresses. If you have less, stretch it to four months rather than cutting steps.

Weeks 1–2 — Find out what's true

No building, no writing, no tools. Just evidence.

- ☐ List every company that has ever paid, trialed, or had a real conversation with you. Everything you know about each.
- ☐ Book six conversations: three customers or best-fit users, two who said no, one who churned.
- ☐ Run the positioning interview script (Chapter 2). Record them. Write down answer 6 verbatim every time.
- ☐ Run the PMF survey if you have 30+ recent active users (Chapter 5).
- ☐ Pull your retention curve, even if it's ugly. Especially if it's ugly.
- ☐ Write down, in one page, what your ten best conversations had in common.

End of week 2 you should be able to say: the three real alternatives to us, the words customers use to describe us, and the characteristics of the people who stick.

Weeks 3–4 — Decide

This is the fortnight where you make the decisions everything else inherits. Do it in long blocks, not in gaps between other work.

- ☐ Complete Worksheet 2 — the positioning canvas. One page, all six boxes.
- ☐ Choose your market frame. It should be niche domination unless you have a very good argument otherwise.
- ☐ Complete Worksheet 3 Part A — the five ICP layers, including the negative ICP.
- ☐ Complete Worksheet 4 — point of view, one-liner, two or three value themes, proof for each, answers to the five objections.
- ☐ Complete Worksheet 7 — value metric, price ladder, discount policy.
- ☐ Run the five-second test with five people in your ICP. Rewrite. Run it again.

End of week 4: five completed worksheets and a one-liner that a stranger repeats back correctly.

Weeks 5–6 — Build the minimum surface

Small, fast, deliberately unfinished. You'll rewrite all of it in six months, and it needs to exist now.

- ☐ Rewrite your homepage to the structure in Chapter 4. Headline, audience, proof above the fold, problem, value themes with proof, how it works, objections, one CTA.
- ☐ Write the one-page internal-sell kit your champion will forward.
- ☐ Define and instrument your activation event (Chapter 8). Find out what percentage reach it and how long it takes.
- ☐ Set up the six-field CRM and the weekly one-pager.
- ☐ Buy and start warming two sending domains. This takes 30 days, which is why it starts now.
- ☐ Set up call recording.

Weeks 5–8 — Build the list, in parallel

Overlaps with the above deliberately. It's the highest-value repetitive work you'll do all quarter.

- ☐ Build 100 scored accounts using Worksheet 3 Part B. At least 25 Tier A.
- ☐ For every Tier A account: named person, verified trigger with a date, and a specific opening angle written out.
- ☐ Pick your two channels using Worksheet 9. For most readers: outbound plus founder-led distribution.
- ☐ Write one channel test card per channel, with a kill criterion, before starting either.
- ☐ Start posting. Three a week, minimum, from week 5, whether or not it feels like it's working. It won't, for a while.

Weeks 9–12 — Run it and read it

The doing phase. Protect the blocks; everything else in the company will try to eat them.

- ☐ Contact all 100 accounts across a 4–6 touch sequence over three weeks. Tier A gets a human-written first line.
- ☐ Target 20+ real conversations in this period. Run the discovery structure every time.
- ☐ Log the quantified pain, in their words, for every conversation.
- ☐ Demo only after discovery. Use their data where you can get it.
- ☐ Any deal past the demo gets a mutual action plan (Worksheet 6).
- ☐ Ask every lost deal, two weeks later, for the real reason.
- ☐ Complete the weekly one-pager every Friday. Every week. No exceptions.
- ☐ Keep posting. Keep the outreach block even in a busy week.

Day 90 — The review

Half a day, alone, with the data in front of you rather than in your head.

FIGURE 15 — THE DAY-90 QUESTIONS

QUESTION	WHAT THE ANSWER CHANGES
What are our actual conversion rates at each step?	Replace the benchmark assumptions in Worksheet 10 with real numbers. Rebuild the model.
Which step is furthest below benchmark?	That's your single focus for the next quarter. Not all of them — that one.
What did every lost deal have in common?	Messaging backlog, product backlog, or ICP correction.
Which accounts converted fastest, and what did they share?	Sharpen the ICP. Usually you'll narrow again, and it'll be right again.
Did either channel hit its success criterion?	Double down on the one that did. Kill the other or run one more cycle with a changed variable.
Can I write the repeatability paragraph yet?	If yes, you're ready for Chapter 13. If not, run another 90 days.

THE REPEATABILITY PARAGRAPH

“[This kind of company], contacted [this way], with [this message], agrees to a conversation about [X]% of the time. About [Y]% of those become customers, in about [Z] weeks, at an average of [£N]. The most common reason we lose is [reason].”

Being able to write that with real numbers is the actual output of the first 90 days. It's also, not coincidentally, what an investor is trying to establish in a first meeting and what a first sales hire needs in order to succeed.

What to do if 90 days produces nothing

It happens, and the instinct — do the same thing harder — is usually wrong. Diagnose in this order, because the cheapest fixes are at the top:

- 1. Nobody replied.** Almost always the list, not the copy. Wrong companies, no trigger, or wrong person. Rebuild 50 accounts with much tighter criteria before you touch a single sentence.
- 2. People replied but wouldn't meet.** Your opening is interesting and your ask is too large. Reduce it — a question, a two-line answer, a specific piece of work you'll do for free.
- 3. Meetings happened but no opportunities.** Positioning or ICP. You're reaching people for whom the problem is real but not urgent. Go back to “who cares a lot.”
- 4. Opportunities but no closes.** Look at what's blocking: no economic buyer (you're selling to users), no quantified pain (you skipped the money question), or price/value mismatch.
- 5. Closes but they churn.** The hardest one, and it's a product or expectation problem, not a marketing problem. Stop selling and go and find out why. Selling more into a leaking bucket is the most expensive mistake in this book.

Notice that four of the five diagnoses send you back to Part I. That's not a coincidence and it's not a failure of the plan. Go-to-market is a loop, and each pass through it sharpens the position, which makes the next pass cheaper.

Scaling Past Yourself

The reward for building a working go-to-market engine is that you become the constraint. Every conversation still routes through you, every deal needs you in the room, and the thing you built to grow the company is now capped by how many hours you have. This chapter is about handing it over without breaking it.

The trigger isn't revenue

Most founders carry sales themselves somewhere between the first \$500K and \$1.5M of ARR, with the median around \$1M.¹³ Those numbers are useful as orientation and they're the wrong thing to decide on.

The real trigger is **repeatability**. A founder at \$2M with chaotic, unexplained conversion rates cannot hire successfully — they'll hand a new rep a process that only works when the founder runs it, and then conclude the rep was bad. A founder at \$500K with clean, understood rates can hire and have it work. The metric beats the milestone.¹³

Three conditions, all of which need to be true:

- You can write the repeatability paragraph from Chapter 12 with real numbers.
- You've closed enough deals the same way that you're bored of them — roughly fifteen to twenty.
- You have more qualified pipeline than you can personally work. Hiring to *create* demand you've never created yourself almost never works.

A fourth, quieter one: you have to be genuinely willing to let someone else do it worse than you for six months. Founders who can't tolerate that hire a rep, override them constantly, and end up with the same workload plus a salary.

Who, first

Not an SDR. This is the most common expensive mistake at this stage.

The instinct is to hire someone junior to do the top of the funnel while you keep the interesting conversations. It fails because a junior person prospecting into a market you haven't fully mapped, with a message you're still tuning, and no ability to run the deal they book, produces meetings you then have to take anyway. You've bought yourself a calendar, not capacity.

Hire a **full-cycle account executive** — someone who can prospect, run discovery, demo, and close, end to end. Three to seven years in B2B software, ideally with experience at your deal size and in something adjacent to your market. Splitting the SDR and AE roles before there's enough pipeline to justify it gives you two half-utilised specialists who convert worse than one person owning the whole thing.⁴¹

FIGURE 16 — HIRING ORDER

STAGE	HIRE	WHY HERE
Pre-\$500K ARR	Nobody	You're still learning. Every conversation you delegate is information you don't get.
\$500K-\$1.5M	First full-cycle AE	Proves the motion transfers to someone who isn't you. This is the real test.
~\$1M+	Second AE, or a GTM engineer	Second AE only after the first hits quota. The systems hire often produces more leverage than another seller.
\$1.5M-\$3M	Marketing or demand generation	Once outbound is proven, someone to build the compounding half.
\$2M-\$4M	Customer success	Earlier if NRR is under 100%. The install base is a pipeline source, not an afterthought.
\$3M-\$5M	Sales leader	Hire a manager to scale an engine that already works, not to build one.

Two caveats. If your motion is genuinely self-serve at low ACV, this order changes — your first hire is more likely a growth or product person than a seller. And a VP of Sales hired before roughly \$1–2M ARR is a well-documented way to spend a year and a lot of money discovering that leaders scale engines rather than invent them.⁴¹

The document you write before you post the job

You cannot hand over a process that exists only in your head, and you'll find out how much of it does when you try to write it down. Budget a full day. This document is also the best possible interview material — give a candidate the ICP and the objection list and ask them to work a deal.

WORKSHEET 12

The sales playbook — contents

- ☐ **Positioning canvas** and the one-liner (Worksheets 2 and 4)
- ☐ **ICP definition**, scoring rubric, and negative ICP (Worksheet 3)
- ☐ **Persona notes** — what each role cares about, what they're measured on
- ☐ **The pitch** — 30 seconds, 2 minutes, 20 minutes
- ☐ **Discovery script** and the four questions that do the work
- ☐ **Demo structure** — the four acts and what to show in each
- ☐ **Qualification criteria** — what must be true to call a deal qualified
- ☐ **Objection handling** — the five, with your actual answers
- ☐ **Pricing and discount policy** — including what they can approve alone
- ☐ **Pipeline stages** and the exit criteria for each
- ☐ **Three recorded calls** — one great, one lost, one that stalled
- ☐ **Case studies** — three, with numbers
- ☐ **Common loss reasons** and what to do about each

The first 90 days of your first hire

- **Weeks 1–2.** They read the playbook, listen to ten recorded calls, and use the product properly. They sit silently on every call you run.
- **Weeks 3–4.** They run discovery, you run the rest. You debrief every call the same day, specifically and unsentimentally.
- **Weeks 5–8.** They own deals end to end with you on the call as support. You speak only when asked.
- **Weeks 9–12.** They're solo, with a weekly pipeline review. First closes should land in this window for short-cycle products; longer for anything above \$50K ACV.

Quota: don't set a full one for the first quarter. Set activity and pipeline-creation targets instead, then move to revenue once you know their ramped rate. For context, across 150+ B2B software companies, about 62% of *ramped* reps hit quota in 2025⁴² — so if your plan assumes 100% attainment from month one, your plan is fiction.

THE HARDEST PART OF THIS TRANSITION

You will lose deals your new hire would not have lost if you'd taken the call. You have to let that happen, or nothing ever transfers. The right response to a lost deal in month two is a debrief and a playbook update, not taking the next one back.

What breaks, and when

FIGURE 17 — PREDICTABLE BREAKAGES

AROUND	WHAT BREAKS	THE FIX
\$250K ARR	The founder is the bottleneck and product stops progressing	Ring-fence time. Two days a week entirely non-sales, defended.
\$500K	Nothing is written down, so nothing can be delegated	Write the playbook. Before you need it.
\$1M	The first hire doesn't sell like you, and win rate drops	Usually a playbook gap, not a person gap. Find the specific stage where they lose and fix that stage.
\$1–2M	Churn becomes visible for the first time as cohorts mature	Somebody owns retention. If NRR is under 100%, that's the whole priority.
\$2–3M	Outbound saturates the narrow ICP you defined	Widen deliberately to an adjacent segment — one at a time, with its own positioning, from strength.
\$3–5M	The founder's personal audience stops being enough on its own	Multi-voice content. The founder anchors, others carry volume. ³¹
\$5M+	Reporting breaks because nobody owns the data	Revenue operations. Earlier than feels comfortable.

The thing not to hand over

Delegate the pipeline. Keep the listening.

The most common pattern in companies that stall between \$2M and \$5M is a founder who stopped talking to customers around the time they hired their second salesperson, and who is now making positioning and roadmap decisions from a dashboard. The dashboard tells you what happened. It never tells you why, and “why” is the only input to the decisions that actually matter.

Five customer conversations a month, forever, regardless of headcount. It's two hours. It is the highest-return two hours in the business and it's the first thing that gets cut.

WHERE THIS ENDS

You started this book with a product and no idea how to sell it. If you've done the work, you now have a written position, a defined audience, a message that lands, a price with a reason behind it, two channels producing conversations, a model that tells you what's broken, and a document that lets someone else run it.

That's a go-to-market engine. It'll need rebuilding roughly every time you double, and each rebuild is faster than the last, because by then you'll know which parts were load-bearing.

The Worksheets

Fourteen worksheets. Twelve appear inside the chapters where they belong; the last two are here because they're for after the book. Print them, or rebuild them in whatever you actually use. What matters is that they get filled in, not where they live.

#	WORKSHEET	CHAPTER	WHAT IT FORCES YOU TO DECIDE
1	The honest diagnostic	1	Which of the six engine parts is actually broken
2	Positioning canvas	2	What you're best at, for whom, against what
3	ICP definition & account list	3	Who exactly, and 100 real companies to prove it
4	Messaging map	4	Your point of view, one-liner, proof and objections
5	PMF scorecard	5	Whether you're allowed to spend money yet
6	Mutual action plan	6	Every step from demo to signature, with dates
7	Pricing decisions	7	Value metric, ladder, discount floor, expansion path
8	Motion selection	8	Self-serve, sales-led or hybrid — and your activation event
9	Channel test card	9	What success and failure look like, before you start
10	GTM model, one page	10	How many conversations a week the target actually requires
11	Automation audit	11	What to hand to a machine and what never to
12	Sales playbook contents	13	What must be written down before you hire
13	90-day tracker	A	Whether you're actually doing the plan
14	Quarterly GTM review	A	What changed, what to stop, what to double

WORKSHEET 13

90-day tracker

One row per week. Fill it in on Friday, in ten minutes, by hand.

WK	NEW CONTACTS	MEETINGS	OPPS CREATED	CLOSED WON	ONE THING I LEARNED ABOUT THE BUYER
----	-----------------	----------	-----------------	---------------	-------------------------------------

1

2

3

4

5

6

7

8

9

10

11

12

The last column is the point. Twelve specific things you didn't know about your buyer in January is worth more than any single quarter's revenue at this stage.

WORKSHEET 14

Quarterly GTM review

1. THE REPEATABILITY PARAGRAPH, UPDATED WITH THIS QUARTER'S REAL NUMBERS

2. CONVERSION RATE BY STAGE — OURS VS LAST QUARTER VS BENCHMARK

STAGE	THIS Q	LAST Q	BENCHMARK
Contact → meeting			3–10%
Meeting → opportunity			~40%
Opportunity → won			28–38%
Trial / POC → paid			~50%
Signup → activated			—
Net revenue retention			>100%

3. THE WEAKEST STAGE — AND THE ONE CHANGE WE’RE MAKING TO IT

4. TOP THREE LOSS REASONS THIS QUARTER

5. WHAT OUR BEST THREE CUSTOMERS THIS QUARTER HAD IN COMMON (ICP CORRECTION?)

6. CHANNEL PERFORMANCE — KEEP, KILL, OR ONE MORE CYCLE

7. WHAT WE’RE STOPPING ENTIRELY

8. THE ONE NUMBER WE’RE MOVING NEXT QUARTER

Benchmarks

Every number in this book, in one place. Read the sample description before you use any of them — most of this data comes from companies substantially larger than yours, and a benchmark applied to the wrong stage is worse than no benchmark at all.

THE PRINCIPAL SOURCE

Most figures below come from ICONIQ’s *The State of Go-to-Market in 2026*, a January 2026 survey of GTM executives — CROs, heads of sales, CEOs and heads of revenue operations — at 150+ B2B software companies, supplemented with quarterly operating data from public and private companies. Respondent revenue: 13% at \$10–25M, 35% at \$25–100M, 27% at \$100–250M, 16% at \$250–500M, 9% above \$500M. By motion: 62% sales-led, 31% hybrid, 7% product-led. By ACV: 7% under \$10K, 20% at \$10–50K, 22% at \$50–100K, 42% at \$100–500K, 10% above \$500K.¹⁵

What that means for you: the median respondent is a company with real scale selling six-figure deals. Their conversion rates are achieved with dedicated teams. Use them for direction and relative comparison — is my demo-to-close far below or near the norm — not as targets to hit in year one.

Funnel

CONVERSION RATES, AVERAGE

STAGE	2025	2026	AI-INFLUENCED PIPELINE >50%
New lead → MQL	25%	28%	38%
MQL → SQL	27%	30%	37%
SQL → closed-won	25%	28%	29%
Demo → closed-won	36%	38%	40%
Free trial / POC → paid	36%	50%	—

WIN RATE BY OPPORTUNITY SOURCE, 2026

SOURCE	2026	2025	NOTE
Customer success	52%	—	Highest. Selling where value is already proven
Sales	38%	43%	Seller-sourced

Channel / partner	35%	39%	Under-used early
Marketing	23%	27%	Volume, lower intent

Cycle, cost and coverage

SALES CYCLE BY ACV, WEEKS, 2026 AVERAGE

ACV	WEEKS	RECOMMENDED PIPELINE COVERAGE
Under \$10K	6	3×
\$10K–\$50K	12	3–4×
\$50K–\$100K	17	4×
\$100K+	24	5–6×

Cycle data: ICONIQ 2026.¹⁶ Coverage: common industry practice; mathematically, required coverage = $1 \div \text{win rate}$.³⁵
Reported average coverage in the ICONIQ dataset was 3.6–3.9×.³⁶

COST PER LEAD AND PER OPPORTUNITY, AVERAGE

SEGMENT	CPL 2025	CPL 2026	CPO 2025	CPO 2026
SMB	\$500	\$500	\$5,200	\$6,300
Mid-market / enterprise	\$800	\$600	\$9,300	\$8,000
Strategic	\$1,300	\$800	\$15,900	\$11,200

Retention, contracts and quota

NET DOLLAR RETENTION, TOP QUARTILE, COMPANIES UNDER \$50M ARR

PERIOD	2023 H1	2023 H2	2024 H1	2024 H2	2025 H1	2025 H2
NDR	118%	124%	115%	113%	121%	123%

CONTRACT LENGTH FOR NEW-LOGO SUBSCRIPTIONS, SHARE OF DEALS

TERM	2023	2024	2025	2026
Under 1 year	4%	4%	2%	13%
1 year	41%	26%	40%	38%
2 years	23%	28%	19%	20%
3 years	28%	36%	34%	23%

4+ years	4%	6%	6%	6%
----------	----	----	----	----

QUOTA ATTAINMENT — % OF RAMPED REPS HITTING QUOTA

CUT	2023	2024	2025
All	59%	58%	62%
By segment, 2025	SMB 68% · mid-market 59% · enterprise 64% · strategic 61%		

Pricing and motion mix

PRIMARY PRICING MODEL — ≥70% OF REVENUE FROM THAT MODEL

MODEL	2025	2026P	PROPRIETARY-MODEL PRODUCTS, 2026P
Hybrid	48%	48%	50%
Subscription / platform	24%	22%	16%
Seat-based	14%	15%	9%
Consumption / usage	9%	9%	16%
Outcome-based	5%	6%	9%

REVENUE SPLIT BY GTM MOTION, AVERAGE

MOTION	ALL, 2026P	HIGH-GROWTH, 2026P	OTHER, 2026P
Direct sales	57%	53%	58%
Channel / partnerships	31%	27%	33%
Self-serve	11%	19%	8%

NEW-LOGO PIPELINE DISTRIBUTION, COMPANIES UNDER \$100M REVENUE

SOURCE	HIGH-GROWTH	OTHER COMPANIES
Sales	62%	47%
Channel / partnerships	19%	15%
Marketing	12%	34%
Other	7%	4%

Self-serve conversion

FREE-TO-PAID CONVERSION — ~200 B2B SOFTWARE PRODUCTS, JANUARY 2026²⁶

MODEL	SHARE OF PRODUCTS	FREE-TO-PAID WITHIN 6 MONTHS
Free trial (all)	57%	Median 8% across all products
— card required	—	~30%
— no card	—	~5%
Freemium	26%	2–8% typical
Reverse trial	7%	—
Interactive demo	7%	—

Typical respondent: \$1–10M ARR, \$50–249 ARPU, 25–50% YoY growth. 14-day trials are the most common length (62%); around 80% of trial-based products include human touchpoints for enterprise users.

AI in GTM

HIGH AI ADOPTERS VS MEDIUM/LOW ADOPTERS, 2026

MEASURE	HIGH ADOPTION	MEDIUM / LOW
Ramped reps hitting quota	67%	59%
SMB / mid-market quota attainment	106% / 91%	80% / 80%
Enterprise / strategic quota attainment	96% / 109%	77% / 88%
Cost per lead / per opportunity	\$670 / \$8.5K	\$600 / \$8.9K
Sales cycle (weeks)	19.3	19.6
Pipeline coverage	3.9×	3.6×
Late renewals	27%	29%
GTM headcount at \$10–25M revenue	20	35
GTM headcount at \$25–100M revenue	45	65

MOST-ADOPTED AI USE CASES IN GTM, % OF COMPANIES, 2026

TOP OF FUNNEL		MIDDLE / POST-SALES	
Lead generation	74%	Call & meeting transcription	75%
Lead scoring & prioritisation	63%	AI for outbound / prospecting	55%

Automated content & campaigns	58%	AI-driven forecasting	38%
Website chatbots & forms	54%	Customer sentiment analysis	45%
AI-powered personalisation	53%	Automated support & ticketing	44%

Glossary

Every term used in this book, defined for someone who has never worked in sales or marketing. Where an industry definition is contested, I've given the useful version rather than the correct-by-committee one.

TERM	PLAIN ENGLISH
ACV	Annual contract value. What one customer pays you in a year. The single number that determines how you're allowed to sell.
Activation	The moment a new user first sees your product do the valuable thing with their own data. The strongest predictor of whether they'll pay.
AEO	Answer engine optimisation. Getting your company named in answers from ChatGPT, Claude, Perplexity and similar — a different discipline from ranking on Google.
ARPU	Average revenue per user. Total revenue divided by number of customers or users.
ARR	Annual recurring revenue. The yearly value of your subscriptions, excluding one-off fees. The default measure of size for software companies.
Bowtie	A model of the customer journey that mirrors the funnel after the sale — onboarding, adoption, renewal, expansion — because in subscription businesses most of the value comes after the close.
Bullseye	A method for choosing marketing channels: brainstorm all of them, rank into three tiers, test three cheaply, then concentrate on the one that works.
Buying committee	Everyone at the customer whose agreement is needed. Averages around eleven people in B2B, most of whom you'll never speak to.
CAC	Customer acquisition cost. Everything you spent on sales and marketing, divided by the customers it produced.
CAC payback	How many months of gross profit it takes to earn back what you spent acquiring a customer. The most useful efficiency metric when cash is tight.
Category entry point	The situation that makes someone start looking for a product like yours — a renewal date, a new hire, a problem that got too big. You want to be the name they think of at that moment.
Champion	The person inside the customer who argues for you when you're not in the room. Usually someone whose own performance improves if you're bought.
Churn	Customers leaving. Logo churn counts customers; revenue churn counts money, which matters more.
Closed-won / closed-lost	A deal that ended in a sale, or didn't.
Cold outreach	Contacting someone who has never heard of you. Legal, effective when targeted, and destructive when done at volume with no relevance.
Consumption pricing	Charging for what gets used — documents processed, API calls, messages sent — rather than a flat fee.

Conversion rate	The percentage of people who move from one stage to the next. Every stage has one, and the weakest one is your priority.
Cost per lead / opportunity	Marketing and sales spend divided by leads or qualified opportunities produced. Cost per opportunity is the more honest of the two.
Demand capture	Reaching people who are already looking for a solution. High intent, small audience.
Demand creation	Making people aware they have a problem worth solving before they start looking. Slow, compounding, hard to attribute.
Deliverability	Whether your emails reach the inbox rather than spam. Governed by domain reputation, authentication and volume discipline.
Discovery call	The first real conversation with a prospect, where you find out what their problem is and what it costs them. Not a pitch.
Economic buyer	Whoever's budget the money comes from. The person who can say yes when everyone else can only say no.
Expansion	Existing customers paying you more — more seats, more usage, more products. Cheaper than new customers and usually neglected.
Firmographics	The structural facts about a company: industry, size, location, funding, business model.
Freemium	A permanently free tier alongside paid tiers. High signup volume, low conversion.
Free trial	Full access for a limited period. Higher conversion than freemium, fewer signups.
Funnel	The sequence of stages from stranger to customer. Useful shorthand, incomplete because it stops at the sale.
Gross margin	Revenue minus the direct cost of delivering it (hosting, model inference, support), as a percentage. Software businesses usually want 70%+.
GTM	Go-to-market. Everything involved in getting a product from "it exists" to "people pay for it."
GTM engineer	Someone who builds go-to-market systems — data pipelines, scoring, signal monitoring, automation — instead of working accounts by hand.
Hybrid motion	Self-serve entry with human sales involvement at the point where deal size justifies it. The mainstream approach in 2026.
ICP	Ideal customer profile. The type of company where you win most and keep them longest. A filter, not a wish list.
Inbound	Prospects finding you and reaching out.
Land and expand	Selling a small initial deal deliberately, then growing the account over time.
Lead	Someone who has shown any interest. Cheap, plentiful, and mostly not worth counting.
LTV	Lifetime value. Total gross profit you expect from a customer before they leave. A projection, and largely meaningless before you have real retention data.
Magic number	New ARR added divided by prior-period sales and marketing spend. A rough measure of whether growth spend is efficient.
MEDDPICC	A checklist for qualifying complex deals: metrics, economic buyer, decision criteria, decision process, paper process, identified pain, champion, competition.

MQL / SQL	Marketing-qualified and sales-qualified lead. Internal labels for how serious someone looks. Useful as measurement, corrosive as a target.
Mutual action plan	A shared document listing every step from here to signature, with owners and dates. The single most useful artefact in a B2B sale.
NRR / NDR	Net revenue (or dollar) retention. What your existing customers are worth this year versus last, including expansion and churn. Above 100% means the base grows on its own.
Opportunity	A qualified deal in progress: real problem, real budget, real timeline. Distinct from a lead.
Outbound	You contacting them first, deliberately.
Paper process	Everything between “yes” and money: legal review, security questionnaires, procurement, purchase orders. Routinely adds weeks nobody planned for.
Persona	The individual role you sell to — their job, targets and daily reality. Distinct from the ICP, which describes the company.
Pipeline	The total value of real opportunities in progress. Not signups, not subscribers.
Pipeline coverage	Open pipeline divided by the revenue you need to close. Mathematically, you need $1 \div \text{win rate}$.
PLG	Product-led growth. The product itself does the acquiring and converting; humans get involved only when the deal justifies it.
PMF	Product-market fit. The state where a definable group of people would be genuinely upset if your product disappeared, and it shows in retention.
POC / pilot	Proof of concept. A time-boxed trial with written success criteria. Converts well when managed, badly when it’s just a login.
Positioning	Deliberately deciding how you’re the best at something a defined market cares a lot about.
Quota	The revenue target a salesperson is expected to close in a period.
Ramp	The time a new salesperson takes to reach full productivity. Usually one to two sales cycles.
Retention curve	The percentage of a cohort still active over time. Whether it flattens is the clearest early signal of product-market fit.
Sales cycle	Time from first real conversation to signature. Scales predictably with deal size.
SDR / AE	Sales development rep (books meetings) and account executive (closes deals). Don’t split these roles until you have enough pipeline for both.
Signal-based outbound	Contacting people because something specific just happened at their company, rather than because they match a filter.
Technographics	What software a company runs. Often a better targeting filter than company size.
Trigger	An event that creates a reason to act now — a hire, a funding round, a threshold crossed, a contract ending.
Value metric	The unit you charge for. Chosen well, your revenue rises as the customer gets more value.
Win rate	Deals won as a share of deals decided (won plus lost). Deals still open don’t count.
95:5 rule	The finding that only around 5% of potential buyers are in a buying process at any moment. The other 95% are why brand and content matter.

Sources

Every external figure in this book, numbered as it appears in the text. Where a number comes from a survey, I've named the sample so you can judge how much weight it deserves at your stage. Where something is an illustration or a rule of thumb rather than measured data, I've said so in the text rather than dressing it up as research.

#	CLAIM	SOURCE
1	Failure and revenue rates for thin application-layer AI products; what constitutes a real moat	Compiled industry analyses of AI application-layer defensibility, 2026 — Baytech Consulting, Preuve, SeedScope, Hatchworks. Directional; sample and methodology vary by source.
2	Average B2B buying committee ~11 people; buyers spend ~17% of purchase time with all suppliers combined; larger groups reduce ability to decide	Gartner B2B buying research, as reported by The Drum, Dock, Madison Logic and Annuitas.
3	The 95:5 rule — ~5% of buyers are in-market at any time; category entry points and mental availability	Prof. John Dawes, Ehrenberg-Bass Institute for Marketing Science, published with the LinkedIn B2B Institute.
4	Generative AI now sits at the top of B2B buyer research interactions, ahead of search and peer referral	Forrester, <i>State of Business Buying 2026</i> , as reported in AEO industry coverage.
5	Only ~12% of pages ranking in Google's top three were cited by major AI assistants for B2B SaaS buyer queries	Independent analysis of 50 B2B SaaS buyer queries, reported by Danish Lead Co / CompetLab, 2026.
6	The “polite call” failure pattern in technical founder sales	Founder-led sales practitioner guides, 2026 — Dowhatmatter, Justin McKelvey, No BS Startup Coach.
7	Definition of positioning	April Dunford, <i>Obviously Awesome: How to Nail Product Positioning so Customers Get It, Buy It, Love It</i> (2019).
8	The ten-step positioning process and its five components	Dunford, as summarised by Heinz Marketing and applied in published case studies (Userlist, Paperless.io).
9	Cold email reply rates ~3–4% vs 15–25% for signal-based outbound; 30–40% lower cost per qualified meeting; ~30 emails per inbox per day; 30-day warmup	2026 outbound benchmark compilations — SpurlQ, Unify GTM, Autobound, Instantly. Vendor-published; treat as directional.
10	The 40% test; benchmarking of 100+ startups	Sean Ellis, as documented by Reforge, Learning Loop and PMF Tracker.
11	Superhuman moving from 22% to 33% by segmentation, then to 58% over three quarters	Rahul Vohra, “How Superhuman Built an Engine to Find Product/Market Fit,” First Round Review.

12	Sub-one-year contracts rising from 4% to 13% of new-logo deals; three-year deals declining	ICONIQ, <i>The State of Go-to-Market in 2026</i> .
13	Founder-led sales typically carried to \$500K–\$1.5M ARR, median ~\$1M; repeatability beats the revenue milestone as the handoff trigger	M Accelerator, No BS Startup Coach, Dock, citing Y Combinator's startup library.
14	MEDDPIC components and adoption in B2B SaaS above \$100K ARR	Apollo, Salesmotion, Weflow, Fullcast, 2026 practitioner guides.
15	Funnel conversion rates 2025 vs 2026; POC discipline; survey firmographics	ICONIQ, <i>The State of Go-to-Market in 2026</i> — January 2026 survey of GTM executives at 150+ B2B software companies, plus quarterly operating data from public and private companies.
16	Sales cycle by ACV: 6 / 12 / 17 / 24 weeks; overall cycles shortened ~6 weeks in 2025	ICONIQ 2026.
17	Win rate by opportunity source — CS 52%, sales 38%, channel 35%, marketing 23%	ICONIQ 2026.
18	Primary pricing model mix, and differences by product technical differentiation	ICONIQ 2026.
19	Prepaid commitment vs pay-as-you-go by revenue scale and ACV	ICONIQ 2026.
20	Van Westendorp price sensitivity meter and its limitations at early stage	Umbrex, Monetizely, First Principles Ventures.
21	AE compensation shifting toward net new recurring revenue and net dollar retention	ICONIQ 2026.
22	Commission clawbacks on consumption deals, typically 3–6 months post-close for \$50K+ ACV	ICONIQ 2026.
23	Cost per lead and per opportunity by segment, 2025 and 2026	ICONIQ 2026.
24	~60% of companies describe acquisition as hybrid; high-growth companies expect ~19% of revenue from self-serve vs ~8% for peers	ICONIQ 2026.
25	Cursor's growth sequence, organic adoption, bottom-up enterprise expansion, and sales leadership timing	GTMnow, "Deconstructing Cursor's Growth Playbook," 2026.
26	Free-to-paid conversion by model; trial vs freemium adoption; card-required effect; trial length; human touchpoints	Growth Unhinged, <i>2026 Free-to-Paid Conversion Report</i> — survey of ~200 B2B software products, January 2026. Typical respondent \$1–10M ARR.
27	Activation's effect on trial-to-paid conversion	2026 SaaS trial benchmark compilations — Growthspree, Userpilot.
28	New-logo pipeline distribution by source for high-growth vs other companies under \$100M revenue	ICONIQ 2026.

29	Demand creation vs demand capture; the limits of software attribution	Chris Walker / Refine Labs, including the Hybrid Attribution Framework.
30	Founder and personal accounts outperforming company pages; 3–5 substantive posts per week	2026 B2B LinkedIn practitioner analyses — HubSpot for Startups, Averi, Foundera.
31	Founder-led social ceasing to be sufficient as the primary motion around \$5–10M ARR	Growthspree, 2026.
32	The Bullseye framework and the 19 traction channels	Gabriel Weinberg & Justin Mares, <i>Traction</i> .
33	The bowtie model of recurring revenue	Winning by Design, <i>The Bowtie Standard</i> .
34	Net dollar retention, top quartile, companies under \$50M ARR, 2023–2025	ICONIQ 2026.
35	Pipeline coverage targets by ACV band; coverage = $1 \div \text{win rate}$	Clari, Forecastio, Landbase, Salesmotion — common industry practice rather than a single study.
36	Reported average pipeline coverage 3.6–3.9x	ICONIQ 2026.
37	CAC payback medians around 15–18 months; efficient companies under 12; PLG motions 6–12	2026 SaaS benchmark compilations — Aleph, Cometly, SaaS Hero. Aggregated secondary data; directional only.
38	High vs medium/low AI adopters: quota attainment, cost per lead and opportunity, cycle length, coverage, late renewals	ICONIQ 2026.
39	Funnel conversion by degree of AI influence on pipeline	ICONIQ 2026.
40	Total GTM headcount by AI adoption level and revenue band	ICONIQ 2026.
41	First sales hire profile; avoiding an early SDR/AE split; VP of Sales timing	First Round Review, RevPilots, Dock, ImpellHub.
42	Ramped rep quota attainment 62% in 2025, and by segment	ICONIQ 2026.

Worth reading in full

- **April Dunford**, *Obviously Awesome* — the positioning book. Short. Read it twice.
- **Rob Fitzpatrick**, *The Mom Test* — how to talk to customers without them lying to you to be nice. Two hours, and it will change every conversation you have.
- **Gabriel Weinberg & Justin Mares**, *Traction* — the nineteen channels and how to choose between them.
- **Winning by Design**, *The Bowtie Standard* — the model behind Chapter 10.
- **ICONIQ Growth**, the annual *State of Go-to-Market* and topline efficiency reports — the best free benchmark data in B2B software.
- **Ehrenberg-Bass Institute** and the **LinkedIn B2B Institute** — if you only ever read one thing about brand, make it their work on category entry points.

ON USING RESEARCH HONESTLY

A number from a survey of 150 companies averaging \$100M+ in revenue is not a target for a three-person company, and anyone who presents it as one is selling something. I've named samples throughout so you can discount appropriately. Where I couldn't verify a figure to a primary source, I've labelled it directional. Your own numbers, once you have twenty data points, outrank everything in this appendix.

The part that actually matters

You can read this book in an afternoon and change nothing. Most people will. The difference between the founders who get to a working engine and the ones who don't isn't intelligence or funding — it's that one group booked six conversations in week one and the other group kept building.

So: three things, this week.

1. **Book three customer conversations.** Use the script on page 12 of Chapter 2. Not a pitch, not a demo. Twenty-five minutes of asking about the past.
2. **Fill in Worksheet 2.** Badly, on one page, in an hour. A rough position you've written down beats a perfect one you haven't.
3. **Start the 90-day plan in Chapter 12.** Put the weeks in your calendar as actual blocks, with actual times.

That's the whole thing. Positioning, then people, then the message, then the conversations, then the numbers that tell you which part is broken. Repeat every time you double.

About the author

I'm Sunny Hooda. I write about marketing and martech as **The Martech Guy**, and I run **Great Hornbills**, where we build websites, run performance marketing and put AI into the boring parts of a go-to-market engine.

Before that, and alongside it: **\$20M in ad spend, managed over ten years. Consumer and B2B.** Which mostly taught me that the ad account is almost never the problem. When a channel isn't working, the cause is upstream nine times out of ten — a position nobody understands, an audience defined too broadly, a message that could belong to anyone. That's why this book spends its first four chapters on things that look like they have nothing to do with pipeline.

I've watched a lot of good products fail to reach the people who needed them. That's an avoidable outcome and it's the reason this exists.

IF YOU WANT TO GO FURTHER

Send me the one page you're least sure about — your positioning canvas, your ICP, or your homepage headline.

I read everything, and if there's an obvious problem I'll tell you what it is. It costs you nothing and it's usually the fastest feedback you'll get on this stuff.

The Martech Guy

Instagram @_themartechguy · YouTube The Martech Guy

Great Hornbills — web, performance marketing, AI automation

From Zero to Pipeline — The Modern GTM Playbook for B2B SaaS

First edition, 2026. © Sunny Hooda. All rights reserved.

This book may not be reproduced or redistributed in whole or in part without permission. The frameworks are yours to use in your own business, freely and without attribution — that's what they're for.

Benchmark data is reproduced with attribution to its original publishers and remains their property. Illustrative examples are invented and labelled as such. Nothing here is financial, legal or investment advice.